
When Does a Speech Model Know to Hand Off? Reading Competence from Hidden States for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A local speech model can keep a conversation going, and a stronger expert can
2 take over the hard questions. Calling the expert on every turn adds cost and waiting
3 time, so the local model needs an early estimate of whether its own answer will fail.
4 Hidden-state probes give such estimates for text models. It is not known whether
5 this signal survives speech and duplex fine-tuning, audio input, and the timing of a
6 native listen/speak protocol. We study these questions on frozen MiniCPM-o 4.5.
7 We use lightweight probes to read failure signals from hidden states and separately
8 calibrate a gate for expert handoff at answer onset. We find that failure signals
9 can transfer from text to speech, and that readout design matters in native audio
10 generation. On our internal benchmark, selective expert handoff raises answer-
11 content accuracy from 40.8% to 62.9%. A separate serving benchmark measures
12 median time to first answer audio of 0.70 s locally, 2.87 s with aggressive handoff,
13 and 6.67 s with always handoff on the internal pool. These are server PCM-ready
14 times; expert-path tails remain a constraint for real-time use.

15 1 Introduction

16 Spoken interaction leaves little time to choose a response [Stivers et al., 2009]. One practical design
17 pairs a local speech model, which keeps the conversation going, with a stronger expert that handles
18 the hard questions. Calling the expert on every turn adds cost and waiting time. Selective escalation
19 needs something else: an early estimate of whether the local answer will fail, made at the onset of a
20 response.

21 Prior work gives two starting points. For text models, hidden-state probes and verbal self-evaluation
22 can predict whether an answer is correct [Kadavath et al., 2022, Azaria and Mitchell, 2023, Kossen
23 et al., 2024], and intermediate layers can carry information that the output does not show [Orgad
24 et al., 2025]. Text routers use such signals to divide queries between models of different cost [Chen
25 et al., 2023, Ong et al., 2024, Varshney et al., 2026]. For speech, native duplex models decide for
26 themselves when to listen and when to speak [Défossez et al., 2024, Cui et al., 2026].

27 Three things remain unknown. First, it is not known whether a failure probe still works when
28 the query pool changes; a high in-distribution AUC can come from pool identity rather than from
29 competence. Second, it is not known whether a probe fit on text input still reads speech-input states
30 after speech and duplex fine-tuning. Third, it is not known where in a streaming protocol the hidden
31 state is actually available for a routing decision. A duplex model already decides when to speak,
32 but this floor-management decision need not reveal whether it knows the answer. It is reasonable
33 to expect that some competence information survives fine-tuning, because the backbone carried it




Figure 1: **The deployed system.** (1) Speech streams in 1 s chunks into a full-duplex local model that decides for itself when to speak. (2) At that speak commit, a mid-network (L22) linear read scores the turn in ~ 30 ms at answer onset; the triggering call may already return the first text unit. (3) Lower-risk turns stay local; higher-risk turns are sent to a cloud expert while the talker stalls. (4) The returned answer is spoken verbatim in the talker’s own voice. Delivered accuracy on the internal test set rises from 40.8% to 62.9% at 52% escalation in native full-duplex sessions.

before. The layer, position, and modality at which it can be read have not been measured on a duplex speech model.

This gap matters in practice. A gate placed at the wrong layer can look good on its calibration pool and then invert on a new one, so the expert is called on the wrong turns. A probe that fails on audio input cannot be used in a speech system at all. And a read that only becomes available after the answer is generated removes the early decision that motivated escalation in the first place.

We study these questions on frozen MiniCPM-o 4.5 [Cui et al., 2026], with raw-backbone and other speech-model controls. Our contributions are three. First, a layer and position audit shows that text-input LOPO math AUC increases from 0.372 at the final layer to 0.931 at layer 22. With audio calibration and audio input, the final-layer probe does not show this failure, so we use the middle layer for its robustness across the tested conditions, not because the output layer never holds competence information. Matched controls link the late-layer degradation to the tested fine-tuned checkpoints but do not identify its cause. At native answer onset, a matched 2×2 layer-by-readout ablation separates depth from aggregation: combining three positions raises L22 mean AUC from 0.595 to 0.695, while its single-state read trails the final layer (Figure 4). Second, a text-calibrated L22 probe transfers to speech-input states at AUC 0.855 on matched TTS queries and 0.760 on 200 query-disjoint human recordings (Figure 5). Cross-model comparisons extend the modality-transfer pattern to MiniCPM-o 2.6 and Qwen2.5-Omni, with weaker transfer in the frozen-backbone Freeze-Omni control. We separate this representation transfer from the extra calibration that native serving needs. Third, a native-protocol benchmark shows that a separately fitted answer-onset gate selects useful expert calls. It raises internal answer-content accuracy by 22.1 points and exceeds the reported random-mixture reference by 6.8 points at the aggressive tier. A separate serving benchmark measures time to first answer audio with real-time input and causal expert prefixes. We distinguish these server PCM-ready measurements from client-audible latency and from the sessions used to evaluate answer-content accuracy.

2 Related Work

Competence readout. Correctness can be estimated from verbal self-evaluation [Kadavath et al., 2022, Lin et al., 2022], hidden representations [Azaria and Mitchell, 2023, Burns et al., 2023, Marks and Tegmark, 2024], and uncertainty across sampled answers [Kuhn et al., 2023, Farquhar et al., 2024]. Intermediate-layer representations can expose information missing from model outputs [Orgad et al., 2025, Mahaut et al., 2024]; Semantic Entropy Probes amortize uncertainty estimation into a hidden-state read [Kossen et al., 2024]. Our question is how readout location, input modality, and serving protocol affect this signal after speech and duplex fine-tuning.

Selective routing. FrugalGPT, HybridLLM, and RouteLLM allocate work across language models [Chen et al., 2023, Ding et al., 2024, Ong et al., 2024]; AutoMix verifies a drafted answer before escalating [Aggarwal et al., 2024]. Pre-generation activation routers and frozen-state success probes are especially close to our method [Varshney et al., 2026, Lugoloobi et al., 2026]. A linear failure probe is thus not itself the novelty. We test its robustness in speech-model representations and use a threshold to expose a tunable expert-call budget. Predicting local failure remains distinct from predicting the benefit of a particular expert.

Native duplex control. Native spoken-dialogue models coordinate listening and speaking [Défossez et al., 2024, Cui et al., 2026], using state or synchronization mechanisms [Wang et al., 2024a, Zhang et al., 2024, Veluri et al., 2024] or frozen-backbone adapters [Wang et al., 2024b]. BayLing-Duplex [Fang et al., 2026] makes dialogue-state decisions within a single autoregressive LLM by interleaving user audio, assistant text, and assistant audio and introducing four state tokens. Starting from GLM-4-Voice, it learns turn-taking and interruption through supervised full-duplex fine-tuning and timing-focused preference optimization. Its evaluation measures spoken-answer quality and interaction timing, providing a concrete example of learned native floor control.

Competence-aware handoff. DuplexOmni, chronological thinking, and MoshiRAG connect speech to reasoning or retrieval [Huang et al., 2026, Wu et al., 2026, Chien et al., 2026]. Our work studies a decision complementary to native floor control: given a duplex-trained checkpoint, can its hidden states predict local-answer failure and guide escalation to an expert under a tunable call budget? We fit a small readout without updating the speech model. This requires regime-specific calibration and does not assume that a learned turn-taking signal predicts answer correctness.

3 Probe-Guided Expert Handoff

3.1 The local-expert loop

Selective model routing balances response quality against the cost of a stronger model [Ong et al., 2024]. This suits our goal: improve answer accuracy under a limited expert-call budget. We adapt the routing decision to a local full-duplex speech model’s transition from listening to speaking (Figure 1). The local model processes incoming speech and decides when to respond; the gate then chooses whether to continue locally or request expert help. On the expert path, an audio uplink transcribes the question, the expert returns a text answer, and the local talker voices the prepared relay text. The request and relay add cost and waiting time. We use local-answer failure risk as a proxy for the value of escalation: benefit still depends on whether the expert repairs the error and the relay preserves its answer.

3.2 A probe for the handoff decision

Linear *probes* test what information a classifier can read from frozen representations [Alain and Bengio, 2016]. Hidden-state classifiers have estimated statement truthfulness [Azaria and Mitchell, 2023] and pre-generation task success [Lugoloobi et al., 2026]. We use this approach to ask whether a speech model’s existing states contain a directly readable failure signal. Reusing these states avoids generating a separate self-evaluation and leaves the speech model unchanged. The linear restriction makes this a test of accessible information; poor performance does not rule out information requiring a nonlinear readout.

Offline, we collect local answers to calibration questions and use a judge to assign correctness labels. Each label is paired with hidden-state features captured at the chosen read point. We fit a logistic probe to predict failure from these features while keeping the speech model frozen. For features $\phi(x)$ from input context x , the fitted weights w, b give the risk score

$$s(x) = \sigma(w^\top \phi(x) + b). \quad (1)$$

At inference, the probe reads the current features without a reference answer or an online judge. Scores above a language-specific threshold τ_ℓ favor escalation; thresholds select nominal expert-call budgets from calibration scores and stay fixed during evaluation (§4.1). Realized call rates can change when the query distribution shifts. An auxiliary dialogue-act head restricts escalation to information requests, filtering turns such as acknowledgments and stop commands (Appendix D).

3.3 Where should the probe read?

Adapting the probe to duplex speech requires choosing its layer, position, and read time. The final layer feeds the output head, but this does not ensure that its failure signal survives a change in query distribution or input modality.

The native feature vector concatenates the latest L22 tail state, the mean of the tail states, and the mean input state, allowing the probe to combine the current boundary state with pooled context. The choice of L22 is informed by the readout analyses in §4.4, which also evaluate the contribution of feature aggregation.

We read these features when the generation call first returns a listen-to-speak transition. This makes the gate available at the local model’s response boundary without waiting for a completed answer. The call can already contain the first generated text unit, so we call the decision point *answer onset*. Further answer chunks may expose additional failure information, at the cost of a later handoff (Appendix D.3). We fit the native probe and thresholds on features and answer labels from this serving regime to match its audio input and generation behavior. This calibration is separate from the representation-transfer study (§4.5), which alone does not provide a zero-shot native gate.

4 Experiments

4.1 Experimental setup

We evaluate frozen MiniCPM-o 4.5 (9B), with gpt-5.5 as the escalation expert, on the fixed 240-query internal test set and four external English speech-QA pools: Speech TriviaQA, Speech Web Questions, and Speech Llama Questions from OpenAudioBench [Li et al., 2025], and SD-QA from VoiceBench [Faisal et al., 2021, Chen et al., 2024]. The native probe is fitted on 5,221 calibration rows after excluding benchmark overlaps. The answer-content results were rerun with this refitted probe; per-language thresholds use the remaining core rows’ out-of-fold scores and stay fixed across pools. Calibration details and model controls are in Appendix A.

Each query has one sampled native session per arm: local-only, conservative, balanced, aggressive, and always-escalate. We report *answer-content accuracy*, scored from generated local text or intended relay text. The random reference is the analytical mixture $A_{\text{mix}}(r) = (1-r)A_{\text{local}} + rA_{\text{always}}$ at the gate’s realized call rate r . Table 1 uses the original query weights; Figure 2’s Internal curve uses the category weights specified in its caption. Both use the same recorded outcomes. External averages weight the four QA pools equally. For content evaluation, the expert receives the final 30 seconds of the prerecorded question. A separate timing experiment replays the same query IDs at real-time speed, supplies only consumed audio to the expert, and enables audio on both paths. It measures signed server TTFA from the scheduled end of input to the first answer PCM block, excluding candidate and stall audio. This run uses a different frozen gate snapshot from the overlap-remediated content fit; its answers are unjudged and do not establish joint accuracy-latency performance (Appendix D.2). Model, fitting, and judge configurations are in Appendix A.1.

4.2 Internal accuracy and timing cost

The upper MiniCPM block of Table 1 reports answer-content accuracy on the fixed 240-query internal test set: 40.8% locally, then 46.7%, 54.2%, and 62.9% for conservative, balanced, and aggressive. Their realized escalation rates are 8.3%, 25.0%, and 51.7% (Appendix D.1). Aggressive exceeds the table’s random reference of 56.1% by 6.8 percentage points. Per-tier internal comparisons are in Table 9. Aggressive recovers 74.6% of the measured local-to-always accuracy gap while using 51.7% expert calls. This measures selection value relative to the recorded endpoints.

Gains vary across the fixed internal strata: factual QA improves by 52.5 points, while math accuracy is unchanged. Of 60 knowledge questions, 43 exceed the expert’s input window. These descriptive results identify an input constraint and task differences, without attributing the accuracy gap to truncation or changing the test mixture (Appendix A.2).

The separate serving run measures when the first answer audio becomes available (Figure 2). Internal TTFA P50 is 0.70 s locally, 0.89 s at conservative, 1.61 s at balanced, 2.87 s at aggressive, and 6.67 s at always. Aggressive has mean 4.59 s and P95 18.14 s; always has mean 8.01 s and P95

Table 1: **Main results: answer-content accuracy (%)**. Both blocks report the internal test set and four external English speech-QA pools; avg is the equal-pool external mean and Δ is its gain over always-local, computed before rounding. MiniCPM uses one recorded native session per query and arm, with fixed calibration thresholds. Matched random mixes its local and measured always arms at each pool’s realized gate rate (Appendix D.1). NVDA uses native-frame replay with cached outcomes; its Internal column is a post-selection diagnostic. Replay details are in Appendix C.

	Internal	TriviaQA	WebQ	Llama Q.	SD-QA	avg	Δ
<i>MiniCPM native sessions (answer-content accuracy)</i>							
always-local	40.8	62.8	52.8	82.4	48.0	61.5	—
+ gate, conservative	46.7	60.4	59.2	78.8	54.0	63.1	+1.6
+ gate, balanced	54.2	70.0	62.4	82.4	65.0	70.0	+8.5
+ gate, aggressive	62.9	88.4	75.2	85.6	82.5	82.9	+21.4
matched random, aggressive	56.1	81.5	68.9	84.1	73.7	77.1	+15.6
always-escalate (measured)	70.4	96.0	79.6	91.6	87.0	88.6	+27.1
<i>NVDA three-layer answer-onset probe (pass-3 native-frame replay)</i>							
always-local	19.6	35.8	37.6	70.8	31.0	43.8	—
+ gate @15%	31.7	48.4	48.8	80.7	42.0	55.0	+11.2
+ gate @30%	44.6	61.0	59.2	85.8	54.0	65.0	+21.2
+ gate @50%	60.4	76.8	69.6	89.7	66.5	75.7	+31.9
matched random @50%	53.5	65.7	60.2	82.0	60.0	67.0	+23.2
always-expert (ceiling)	91.7	95.5	82.8	93.1	89.0	90.1	+46.3

25.86 s. Local mean is -0.61 s because 17 of 239 completed sessions begin answering before the input ends. Such early responses do not establish successful turn-taking. Across 5,950 attempts, 5,904 complete and 46 fail; distributions condition on completion and retain outliers (Table 10). These measurements characterize server response time, while correctness under this causal input protocol and client-audible delay remain unmeasured.

4.3 External transfer and its limits

The native gate’s weights, feature definition, layer, and per-language thresholds stay fixed across the four external English QA pools. The upper block’s avg excludes the internal set; Δ is its gain over always-local. Aggressive macro accuracy rises from **61.5% to 82.9%**, compared with 77.1% for the reported random mixture: +21.4 points over local answering and +5.9 over random. Its realized call rate ranges from 18.8% on Llama Questions to 66.0% on SD-QA, so this is not a uniform 50%-budget comparison. TriviaQA and SD-QA illustrate large improvements over local answering (+25.6 and +34.5 points) and over the corresponding random reference (+6.9 and +8.8 points).

The lower tiers are less reliable. Conservative accuracy falls by 2.4 points on TriviaQA and 3.6 on Llama Questions; balanced leaves Llama Questions unchanged. Separately sampled sessions and limited expert headroom constrain the interpretation of small differences. The results support aggressive-tier gains across these pools, not improvement on every benchmark at every budget. On AlpacaEval, we measure 3.68 locally and 4.01 at aggressive, with a random-mixture reference of approximately 3.9; the always arm scores 4.20 versus 4.95 for returned expert text. The relay cap is particularly restrictive for long, open-ended answers.

Failure categories and later reads help explain the boundary. In the native failure-type audit, specific-but-wrong answers comprise 70% of failures and are ranked at AUC 0.81, while execution slips in multi-step reasoning are near chance (0.47). In a separate read-time sweep, external mean AUC increases from 0.755 at commitment to 0.791 after two answer chunks. This is consistent with some useful information emerging during generation, without proving a causal mechanism (Appendix D.3).

The lower block of Table 1 reports the exploratory NemotronLabs-VoiceChat-11B (NVDA) study. Its avg covers four English QA pools and excludes the Internal column. At a 50% replay budget, average accuracy rises from 43.8% to 75.7% ($\Delta = 31.9$ points), compared with 67.0% for random. This is native-frame replay with cached outcome mixtures, not another set of MiniCPM native ses-




Figure 2: **Answer-content accuracy and separately measured serving time.** Left: accuracy against expert call rate. Internal uses post-hoc weights of 0.25 for knowledge/math and 1 otherwise; all 240 queries remain included. Table 1 retains the original weights; external curves are unweighted. The dashed line is the random endpoint mixture; the horizontal line is the measured always accuracy. Right: mean and P50 of signed server PCM-ready TTFA from a separate serving run on the same query IDs, with original query weights and a different frozen gate snapshot (Appendix D.2). Statistics use completed sessions; negative times retain early responses. The panels do not establish joint accuracy–latency performance. L/C/B/A/W denote local/conservative/balanced/aggressive/always. Table 10 gives counts and tails for all five pools.

sions. The L26/L30/L34 heads reach 0.818 out-of-fold AUC and 0.816 mean external AUC in the separate probe analysis (Appendix C); Table 1 reports accuracy rather than AUC. External pools were inspected during variant review, so independent confirmation is needed.

4.4 Readout ablations

We evaluate the readout independently of the expert path using two protocols: a text-input layer-and-position sweep and a native-audio layer-by-readout ablation.

Text-input transfer. We test whether late-layer degradation persists across read positions and raw-backbone controls. Under leave-one-pool-out (LOPO) evaluation, probes are fitted on non-math pools and tested on held-out math (Figure 3). A last-token probe at L22 reaches **0.931** AUC on held-out math, compared with the final-layer baseline of 0.372, while mean pooling reduces the late-layer degradation. A raw Qwen backbone, subsampled to match per-pool failure rates, remains near 0.96;



Figure 3: **Text-input transfer depends on layer and position.** LOPO math AUC across layers: late last-token reads deteriorate on the duplex checkpoints, while mid-layer reads and raw-backbone controls remain predictive. Layer 22 is the MiniCPM-o 4.5 read location used for the native gate. This plot does not show an equivalent collapse under audio calibration and audio input.




Figure 4: **Aggregation improves failure prediction at L22.** Mean AUC across five external pools under leave-one-pool-out evaluation on one native audio batch. Gray bars read the latest tail state; blue bars concatenate that state with the tail mean and input mean (§3.3). Arrows show the AUC change from aggregation; chance is 0.5.

207 the Qwen2.5 control family shows a similar ordering (Table 5). These comparisons associate the
 208 degradation with the tested fine-tuned checkpoints, without isolating a causal training objective.

209 With audio calibration and audio input, the final-layer probe reaches 0.936 on LOPO math. The ob-
 210 served late-layer degradation is therefore conditional on the input and calibration setting; its mecha-
 211 nism remains open (Appendix B.2).

212 **Native layer and aggregation.** We test whether predictive gains at native answer onset come from
 213 layer choice or feature aggregation. A 2×2 comparison crosses L22 versus the final layer with a
 214 single tail state versus the three-part aggregate defined in §3.3. All four readouts use one audio batch
 215 with identical training pools, labels, splits, and head recipe (Figure 4). At L22, aggregation raises
 216 mean failure-prediction AUC from 0.595 to 0.695; at the final layer, the change is smaller (0.654
 217 to 0.662). The single-state L22 read has no advantage over the final layer. These results identify
 218 an aggregation benefit at L22, without establishing an advantage from middle-layer selection alone.
 219 Named per-pool results and cached routing gains are reported in Appendix B.4.

220 4.5 Text-trained probes transfer to unseen human speech

221 We next test whether a probe fitted to text-input hidden states can predict failure under speech input.
 222 This tests a change of input modality; Figure 3 tests held-out task-pool transfer within text input.




Figure 5: **Text-trained failure probes transfer to unseen human speech.** Left: matched-query modality transfer—probes fitted on text-input states are evaluated on TTS-audio states for the same 600 questions, across four models. MiniCPM-o 4.5 reaches 0.855 AUC at L22. Right: query-disjoint transfer—MiniCPM-o 4.5 probes fitted on the calibration pool are evaluated on 200 held-out SD-QA human recordings. The text-trained probe (blue) reaches 0.760 AUC at L22 and peaks at 0.800 at L25; gray curves show audio-trained and text-target references. These tests evaluate cross-modal transfer, complementing the text-only task-pool transfer in Figure 3. Depth is normalized by each model’s layer count.

Figure 5 (left) holds question content fixed: the same 600 queries are presented as text and as TTS audio. The text-trained MiniCPM-o 4.5 probe reaches **0.855** AUC on audio states at L22. Transfer also appears in MiniCPM-o 2.6 and Qwen2.5-Omni, with plateau AUCs of 0.74–0.80 and 0.80–0.83, respectively; the frozen-backbone Freeze-Omni control shows much weaker transfer under the tested recipe. These matched-query tests measure modality transfer with question content shared across fitting and evaluation.

The right panel tests generalization to *new questions and human speech*: probes fitted on the calibration pool are evaluated on 200 query-disjoint SD-QA recordings. The text-trained MiniCPM-o 4.5 probe reaches **0.760** AUC at L22, with a per-layer peak of 0.800 at L25. Audio-trained probes and evaluation on SD-QA text provide references for the modality change. Thus, transfer extends beyond matched TTS questions to held-out human speech. The two panels use different evaluation sets, so their AUC difference does not isolate a human-speech penalty (Appendix B.1). These results concern representation transfer; native escalation uses a separately fitted probe with three L22 positions (§3.3).

5 Discussion and Limitations

The central finding is that a mid-layer read remains predictive in tested conditions where a final-layer, last-token probe fails. This read location also supports selective expert calls without updating the speech model.

Three boundaries define the scope of the results. First, coverage is limited to the tested checkpoints and predominantly single-turn speech QA. Broader conversational tasks, including multi-turn dialogue and expert continuity, remain to be evaluated. Second, native serving requires in-regime audio features and answer labels to fit the readout and set thresholds; representation transfer alone does not provide a zero-shot gate. Third, answer-content accuracy and server TTFA come from separate runs with different gate snapshots. The timing run has unjudged answers, early responses, and long tails; it does not establish joint accuracy–latency gains for the overlap-remediated gate. Human-rated speech quality and client playback timestamps remain necessary to characterize the listener’s experience.

These results motivate competence-aware handoff as a useful addition to a speech model’s floor control. Appendix A provides the model and calibration details for the benchmark.

References

- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.
- Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. *arXiv preprint arXiv:1610.01644*, 2016.
- Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2212.03827.
- Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.
- Chung-Ming Chien, Manu Orsini, Eugene Kharitonov, Neil Zeghidour, Karen Livescu, and Alexandre Défossez. MoshiRAG: Asynchronous knowledge retrieval for full-duplex speech language models. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2604.12928.
- Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiyue Sun, Yingjing Xu, et al. MiniCPM-o 4.5: Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.
- Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*, 2024.
- Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2404.14618.
- Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA: Spoken dialectal question answering for the real world. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.
- Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630:625–630, 2024.
- Muye Huang, Lingling Zhang, Xingyu Yu, Lei Shi, Zhanyu Ma, Jun Xu, Jiuchong Gao, Jinghua Hao, Renqing He, and Jun Liu. DuplexOmni: Real-time listening, seeing, thinking, and speaking for full-duplex interaction. *arXiv preprint arXiv:2606.09186*, 2026.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. Semantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint arXiv:2406.15927*, 2024.
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2302.09664.

299 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
300 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
301 Introduces OpenAudioBench.

302 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
303 words. *Transactions on Machine Learning Research*, 2022. arXiv:2205.14334.

304 William Lugoloobi, Thomas Foster, William Bankes, and Chris Russell. LLMs encode their failures:
305 Predicting success from pre-generation activations. *arXiv preprint arXiv:2602.09924*, 2026.

306 Matéo Mahaut, Laura Aina, Paula Czarnowska, Momchil Hardalov, Thomas Müller, and Lluís
307 Màrquez. Factual confidence of LLMs: on reliability and robustness of current estimators. In
308 *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.

309 Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language
310 model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
311 arXiv:2310.06824.

312 Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez,
313 M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data.
314 *arXiv preprint arXiv:2406.18665*, 2024.

315 Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and
316 Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of
317 LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025.
318 arXiv:2410.02707.

319 Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann,
320 Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C. Levin-
321 son. Universals and cultural variation in turn-taking in conversation. *Proceedings of the National*
322 *Academy of Sciences*, 106(26):10587–10592, 2009.

323 Tanay Varshney, Annie Surla, Michelle Xu, Gomathy Venkata Krishnan, Maximilian Jeblick, David
324 Austin, Neal Vaidya, and Davide Onofrio. LLM router: Rethinking routing with prefill activations.
325 *arXiv preprint arXiv:2603.20895*, 2026.

326 Bandhav Veluri, Benjamin N. Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. Be-
327 yond turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In *Proceedings*
328 *of EMNLP (Main)*, 2024. arXiv:2409.15594.

329 Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. A full-duplex
330 speech dialogue scheme based on large language models. In *Advances in Neural Information*
331 *Processing Systems (NeurIPS)*, 2024a. arXiv:2405.19487.

332 Xiong Wang, Yangze Li, Chaoyou Fu, Yunhang Shen, Lei Xie, Ke Li, Xing Sun, and Long Ma.
333 Freeze-Omni: A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv*
334 *preprint arXiv:2411.00774v5*, 2024b.

335 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
336 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
337 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest*
338 *Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

339 Xinrong Zhang, Yingfa Chen, Shengding Hu, Xu Han, Zihang Xu, Yuanwei Xu, Weilin Zhao,
340 Maosong Sun, and Zhiyuan Liu. Beyond the turn-based game: Enabling real-time conversations
341 with duplex models. In *Proceedings of EMNLP*, 2024. arXiv:2406.15718.

A Benchmark and calibration details

The main table reports the original internal query mixture and four external English QA pools. The figure’s Internal accuracy and call-rate curve uses post-hoc knowledge/math weights of 0.25 and weights of 1 for other categories; timing retains the original mixture. Archived Mandarin diagnostics are separate from the main-table macro average.

A.1 Experimental setup details

Models and decoding. The target is MiniCPM-o 4.5 (9B, based on Qwen3-8B). Matched controls include raw backbones, a second MiniCPM generation, Qwen2.5-Omni, Qwen2-Audio, and Freeze-Omni (Table 3). Representation analyses use bf16, greedy decoding, and seed 42. Native decoding uses temperature $T = 0.7$ and top- $k = 20$. The escalation expert is gpt-5.5 with low reasoning effort.

Probe fitting and thresholds. Representation probes use 360 calibration queries and single-position 4,096-dimensional features. The native gate uses 5,221 calibration rows and three L22 summaries: the last tail state, the tail mean, and the input mean. Its logistic head uses $C = 3 \times 10^{-4}$. Five-fold stratified cross-validation uses shuffle seed 42. For each language, conservative, balanced, and aggressive thresholds are the $(1 - r)$ quantiles of the 4,979 core rows’ out-of-fold scores, with $r \in \{0.15, 0.30, 0.50\}$. The 242 fresh training rows enter the fit but not the quantile base.

Judge configuration. Speech TriviaQA, Speech Web Questions, and Speech Llama Questions use the OpenAudioBench judge with gpt-4o-2024-08-06. Internal, SD-QA, and Mandarin reasoning answers use the reference-anchored gpt-5.4-mini adequacy judge. AlpacaEval uses VoiceBench’s gpt-4o-mini judge on its 1–5 scale. The adequacy judge receives the question, an optional reference, and a candidate answer without the system identity. Its rubric requires factual correctness, valid reasoning, and a complete answer. Prompt details are in Appendix E.

A.2 Internal task strata

Table 2: Descriptive strata of the fixed internal test set. Local, aggressive (A), and always report answer-content accuracy (%); call rate is the aggressive arm’s realized rate. The final column counts question waveforms longer than the expert’s 30-second input window. One recorded run per query and arm; no samples are removed or reweighted.

Stratum	n	Local	A	Always	A call rate	$> 30\text{ s } (n)$
Chat	60	65.0	70.0	70.0	35.0	0
Factual QA	40	32.5	85.0	92.5	62.5	0
Knowledge	60	15.0	31.7	38.3	71.7	43
Math	60	61.7	61.7	78.3	25.0	3
SimpleQA traps	20	0.0	95.0	100.0	100.0	0

The largest gains occur on factual QA and SimpleQA traps. Math is unchanged at the aggregate level, consistent with the failure-type analysis in Appendix D.3. Long-input truncation is concentrated in the knowledge stratum (43 of 60 questions); all queries remain in the reported denominator.

B Models and representation analyses

B.1 Signal comparison and layer transfer

The final-layer state predicts source pool at 95.8% accuracy, and a pool-identity oracle alone reaches failure AUC 0.715. This motivates LOPO rather than an in-distribution comparison. The depth sweep shows that the failure is localized: mid-layer last-token reads transfer while late reads degrade on both duplex models; mean pooling degrades less. In the signal table, the final-layer audio entry is audio-calibrated LOPO math, whereas the L22 audio entry tests a text-calibrated probe on speech states.

Table 3: Model roster. “Pair” identifies a fine-tuned model and its raw backbone control.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control

Table 4: Failure-prediction AUC on MiniCPM-o 4.5. In-dist uses five-fold out-of-fold predictions; LOPO holds out the named pool.

signal	in-dist	LOPO math	audio input
max entropy@4	0.696	near chance	—
$p(\text{True})\text{-pre}$	0.807	trap 0.945	degrades
$p(\text{True})\text{-post}$	0.899	—	degrades
final-layer probe	0.822	0.372	0.936
mid-layer (L22) probe	0.866	0.931	0.855

377 The text-to-audio result is not explained by synthetic speech alone: on 200 human SD-QA record-
378 ings, the text-calibrated read remains at 0.76–0.80 AUC across L22–L26. Audio-input verbal self-
379 evaluation is less reliable, whereas adding a transcript restores it; the hidden-state probe does not
380 require that extra generation step.

381 B.2 Mechanism checks

382 Matched controls narrow the late-layer effect without assigning it to a specific training objective.
383 A speak-mode template leaves final-layer LOPO math unchanged (0.362 versus 0.366), the audio-
384 text direction is nearly orthogonal to the failing probe, and source-pool identity is decodable at
385 all depths. A residual-component analysis shows the distributed read holding through the middle
386 layers and decaying only near the outputs of the two duplex models; it remains intact in their raw
387 backbones and the omni-streaming control. This supports the practical choice of a mid-layer read
388 while leaving the causal mechanism open.

389 B.3 Query-surface baseline

390 A TF-IDF logistic router reaches 0.669 out-of-fold AUC and 0.38–0.57 under LOPO, including
391 0.230 on the all-failure trap pool. The hidden-state probe’s advantage is therefore not explained by
392 query wording or source identity alone. On audio labels the surface router improves to 0.814 but
393 still trails the L22 probe at 0.879.

394 B.4 Native feature ablation

395 Figure 4 uses one native audio batch (agg0, official configuration), with all layer features captured
396 in the same sessions and judged against their own generated answers. Single-state and aggregated
397 features have 4,096 and 12,288 dimensions, respectively. Each readout fits a logistic head ($C =$
398 3×10^{-4}) on four pools and evaluates the fifth. The five pools include Mandarin reasoning; this
399 set differs from the main table’s four English external pools. Table 6 identifies the individual pool
400 results that underlie the averages, including the weaker reasoning scores.

401 The routing comparison selects the highest-risk 30% of questions in each pool and combines
402 recorded local outcomes with cached always-arm expert outcomes. Figure 6 reports pooled ac-
403 curacy gains over 2,000 matched random allocations. Per-pool gains in Table 6 use the analytical
404 30% random mixture. These are cached outcome comparisons, not new live gate sessions. For the
405 L22 single-state read, 0.0505 of the random allocations matched or exceeded the observed accuracy;

Table 5: LOPO math AUC by depth for matched model families.

model	best mid-layer	final layer	shape
Qwen3-8B	0.964 (L33/36)	0.958	plateau
MiniCPM-o 4.5	0.931 (L22/36)	0.366	late inversion
Qwen2.5-7B	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni	0.794 (L16/28)	0.746	diffuse decline
MiniCPM-o 2.6	0.822 (L21/28)	0.540	late cliff

406 none did for the other three readouts. These comparisons test against random routing, not differences
407 between the readouts.




Figure 6: **Cached routing gains at a matched 30% budget.** Accuracy gains over random escalation, pooled across the same five evaluation sets as Figure 4. Each bar uses the recorded local and expert outcomes; feature choices and colors match the main figure.

Table 6: Native ablation by evaluation pool. Each column is one of the four matched LOPO readouts in Figure 4. The lower block gives accuracy gains over random escalation in percentage points. AUC averages weight pools equally; pooled routing gains weight individual questions.

Pool	L22		L35	
	Single	Aggregate	Single	Aggregate
<i>Failure-prediction AUC</i>				
Speech TriviaQA	0.645	0.744	0.704	0.661
Speech Web Questions	0.579	0.732	0.646	0.683
Speech Llama Questions	0.614	0.711	0.672	0.674
SD-QA	0.648	0.713	0.682	0.669
Reasoning QA (zh)	0.488	0.576	0.564	0.624
Equal-pool mean	0.595	0.695	0.654	0.662
<i>Accuracy gain over random at a 30% budget (pp)</i>				
Speech TriviaQA	+3.38	+4.58	+3.78	+3.38
Speech Web Questions	+1.69	+1.69	+4.16	+3.33
Speech Llama Questions	+1.51	+3.96	+2.73	+3.96
SD-QA	+2.85	+5.35	+5.85	+4.35
Reasoning QA (zh)	-3.91	-1.93	+1.04	+1.53
Pooled gain	+1.20	+2.78	+3.47	+3.33

408 For context, the separately trained deployed gate reaches mean AUC 0.750 and a pooled gain of 4.80
409 percentage points on these rows. It was fitted on the full calibration merge, so it is not a matched
410 training condition for the four LOPO ablations and is excluded from their plots.

C Second duplex family

We replay NemotronLabs-VoiceChat-11B through its native frame protocol with the checkpoint frozen. Of 2,550 English calibration queries, 2,481 produce a commit frame. Standardized logistic heads read L26/L30/L34 onset states; their normalized logits are averaged. The recipe is chosen on calibration-only LOPO metrics, and external labels are excluded from fitting. Each head uses $C = 10^{-4}$ and concatenates the commit state, eighth post-commit state, first-eight-state mean, and running pre-commit mean (17,920 dimensions per layer). External pools were inspected during variant review, so these results provide exploratory transfer evidence. The read covers the first 640 ms after commit, so this experiment tests method transfer rather than end-to-end real-time serving.

Table 7: VoiceChat-11B transfer. Mean accuracy equally weights four English speech-QA pools; random is the matched-rate mixture.

expert budget	15%	30%	50%
probe-guided accuracy	55.0	65.0	75.7
random-mixture accuracy	50.7	57.7	67.0

The three-layer read reaches 0.818 out-of-fold AUC and 0.816 mean external AUC under the pool-aligned judges. A strictly causal commit-only variant loses about 0.03 external AUC, making answer onset the closest matching read point across the two model families.

D Native gate and serving details

The L22 hidden state is available after roughly 57–61% of prefill. On one H100, probe availability has P50/P95 of 20/25 ms for text and 45/104 ms for audio. These numbers measure the local read, not the network or expert call. The native benchmark’s realized rates are below; their variation is expected because thresholds are calibration quantiles rather than per-pool quotas.

Dialogue-act filtering. A second linear head uses the same L22 read to distinguish information requests from floor-management turns. Escalation requires both the information-seeking score $a(x)$ and failure score $s(x)$ to exceed their respective thresholds:

$$d(x) = \mathbf{1}\{s(x) \geq \tau_\ell\} \mathbf{1}\{a(x) \geq \tau_a\}. \quad (2)$$

On the scripted calibration stimuli, the act head has OOF AUC 1.000; the question-side 0.5th-percentile threshold removes 0.5% of true escalations and admits no floor-management stimuli.

433 D.1 Realized escalation rates

Table 8: Realized escalation rates (%) for the MiniCPM native results in the upper block of Table 1, plus the archived Mandarin and separate AlpacaEval pools. These are measured rates, distinct from the nominal calibration budgets and from the NVDA replay budgets. Local-only makes no expert calls.

Pool	n	Conservative	Balanced	Aggressive	Always
Internal	240	8.3	25.0	51.7	99.2
Speech TriviaQA	250	4.0	18.4	56.4	100.0
Speech WebQ	250	4.8	22.0	60.0	99.2
Speech Llama Q.	250	0.4	3.6	18.8	100.0
SD-QA	200	6.5	27.0	66.0	100.0
Reasoning QA (zh)	202	19.8	33.2	50.5	100.0
AlpacaEval	199	1.5	4.5	43.7	100.0

Table 9: Internal native-protocol answer-content accuracy with query-bootstrap 95% confidence intervals ($n = 240$). Random follows Table 1’s endpoint-mixture convention.

tier	realized calls	accuracy	random reference
local-only	0.0%	.408 [.35,.47]	.408
conservative	8.3%	.467 [.40,.53]	.433
balanced	25.0%	.542 [.48,.60]	.482
aggressive	51.7%	.629 [.57,.69]	.561
always	99.2%	.704 [.65,.76]	.704

434 D.2 Measured time to first answer audio

435 **Clock and endpoints.** The serving experiment uses the same 1,190 query IDs across five pools and
 436 five arms, with one attempt per query and arm. Input chunks become available at their scheduled
 437 arrival times under real-time replay; consumption cannot precede arrival. A worker’s monotonic
 438 clock measures

$$\text{TTF}_{\text{server}} = t_{\text{first answer PCM}} - t_{\text{input end}},$$

439 where input end is the scheduled arrival of the last real audio sample. On the local path, the end-
 440 point is the return of the first non-listen chunk containing nonempty PCM. This interface is chunk-
 441 blocking. On the expert path, the endpoint is the first nonempty PCM yielded by the teacher-forced
 442 talker, before whole-waveform synthesis completes. The onset chunk’s candidate audio and subse-
 443 quent stall audio are logged separately and excluded from answer TTF. Negative times are retained:
 444 the model may answer or finish before all input arrives. The measurement does not include network
 445 delivery or client playback and does not certify answer correctness or full-input coverage.

446 **Serving configuration and calibration scope.** Audio generation is enabled on both paths, with
 447 top- $k = 20$, three forced listen chunks, and an L22 read aggregating eight tail states and the input
 448 mean. At escalation, ASR receives only the consumed real audio prefix; gpt-5.5 uses low reason-
 449 ing effort, with web-enabled requests and a non-web fallback. Prepared expert text is relayed by the
 450 teacher-forced talker. After the local turn, expert completion is polled at 50 ms intervals with a 150 s
 451 wait limit. Each worker warms up with six chunks and one synthesis before processing sessions
 452 serially.

453 The frozen timing gate was fitted on 5,228 rows, with English thresholds 0.9057/0.7973/0.5534. It
 454 differs from the 5,221-row overlap-remediated fit described for content evaluation in Appendix A.1.
 455 The seven external-question overlaps remain in this timing gate’s calibration. These timings there-
 456 fore characterize this serving configuration, not the generalization or joint accuracy-latency per-
 457 formance of the overlap-remediated gate. No correctness judgments were collected for the timing
 458 sessions. Hardware logs contain 5,719 sessions on H100 80GB HBM3, 56 on H100 NVL, and 175
 459 on H200; all 1,200 Internal sessions use H100 80GB HBM3. External timings include this hardware
 460 variation and are not controlled hardware comparisons.

Coverage, failures, and tails. All 5,950 attempts are retained: 5,904 completed sessions, 44 without a speak commit, one empty response, and one expert timeout. Smoke runs are excluded. Table 10 conditions timing statistics on completed sessions with a recorded answer endpoint and reports failed attempts separately. Calls is the fraction of all attempted sessions that fired; this denominator includes sessions without onset. No timestamp is zero-filled and no negative value or outlier is clipped. Internal local TTFA has mean -0.61 s but P50 0.70 s, with 17 early responses. Speech Web Questions has long delayed-onset outliers: local TTFA reaches 305.52 s despite P95 3.60 s. Quantiles describe variation across queries, not uncertainty over repeated runs.

On the 235 Internal IDs valid in all five arms, paired-subset P50 values are $0.69/0.88/1.61/2.88/6.64$ s for local/conservative/balanced/aggressive/ always. The corresponding common-ID counts for Speech TriviaQA, Speech Web Questions, Speech Llama Questions, and SD-QA are $248/237/242/199$. The release records per-session timestamps, status, input identity, expert-prefix identity, audio metadata, and frozen configuration hashes.

Table 10: Measured server PCM-ready TTFA (seconds). Each arm attempts every query in the pool. Statistics use completed sessions; Fail counts all other outcomes, and Early counts negative TTFA among completed sessions. Calls uses all attempted sessions as its denominator. One run per query and arm; quantiles describe query variation.

arm	Valid	Fail	Early	Calls (%)	Mean	P50	P95	P99
<i>Internal (n = 240 per arm)</i>								
local	239	1	17	0.0	-0.61	0.70	2.62	3.64
conservative	238	2	12	9.6	1.53	0.89	10.85	34.08
balanced	236	4	11	25.4	2.36	1.61	14.41	26.91
aggressive	236	4	13	50.8	4.59	2.87	18.14	54.78
always	237	3	7	99.2	8.01	6.67	25.86	60.83
<i>Speech TriviaQA (n = 250 per arm)</i>								
local	250	0	1	0.0	1.14	0.90	1.92	2.71
conservative	248	2	0	3.6	1.55	0.98	2.63	14.88
balanced	250	0	0	18.4	2.45	1.46	8.46	23.74
aggressive	250	0	0	56.4	3.62	3.55	7.92	16.07
always	250	0	0	100.0	5.71	4.89	10.30	21.76
<i>Speech Web Questions (n = 250 per arm)</i>								
local	246	4	0	0.0	3.69	1.71	3.60	4.25
conservative	246	4	0	5.6	5.14	1.81	7.71	80.26
balanced	244	6	0	21.2	3.62	2.18	9.88	20.61
aggressive	247	3	0	59.6	7.40	5.07	13.25	24.50
always	248	2	0	99.2	7.49	6.13	16.71	23.13
<i>Speech Llama Questions (n = 250 per arm)</i>								
local	249	1	0	0.0	1.61	1.52	2.55	3.66
conservative	249	1	0	0.0	1.63	1.54	2.55	3.71
balanced	247	3	0	4.0	2.10	1.53	3.82	22.78
aggressive	246	4	0	19.2	2.61	1.68	8.49	13.21
always	249	1	0	99.6	5.13	4.32	9.45	15.69
<i>SD-QA (n = 200 per arm)</i>								
local	200	0	1	0.0	0.93	0.62	1.67	2.24
conservative	199	1	2	6.5	1.50	0.75	7.18	11.94
balanced	200	0	0	24.5	2.60	0.87	10.81	15.52
aggressive	200	0	0	65.0	4.66	4.20	10.17	14.52
always	200	0	0	100.0	5.85	5.10	10.64	14.00

D.3 Failure type and later reads

This separate sweep fits a probe at each read point using 5,236 judged calibration rows, row-random five-fold fitting, and $C = 3 \times 10^{-4}$. Later reads pad finished answers with the last available vector. Its five-pool macro includes Mandarin reasoning, unlike the main table’s four-English-pool macro. Failure categories are model judgments.

Failure-prediction AUC at successive native read points. External macro equally weights five speech-QA pools.

	Read point	Calibration OOF	Internal	Reasoning-zh	External macro
479	Before onset	.848	.832	.711	.753
	At onset	.850	.852	.648	.755
	After 1 chunk	.852	.854	.681	.762
	After 2 chunks	.855	.846	.769	.791
	After 3 chunks	.857	.858	.744	.801

480 Specific-but-wrong answers account for about 70% of failures and are ranked at AUC 0.813; per-
 481 ception errors reach 0.889, while execution slips are near chance at 0.469. The later-read gains on
 482 reasoning expose a tradeoff between earlier routing and information that emerges during generation.

483 **E Judge prompt and reproducibility**

484 The answer judge receives the question, an optional reference, and a candidate answer without
 485 the system identity. It marks the candidate adequate only when it is factually correct, uses valid
 486 reasoning, and completely answers the question; partial, hedged, refused, off-topic, and truncated
 487 answers are inadequate. The pre-answer self-evaluation asks the model not to answer the question
 488 and to reply exactly “Yes” or “No” to whether it would answer correctly. The post-draft version
 489 presents the question and proposed answer and asks whether the answer is correct.

490 Representation-study sampling and splits use seed 42. The native benchmark uses sampled decoding
 491 and retains one run per query and arm. MiniCPM-o 4.5 runs with PyTorch 2.8.0 and Transformers
 492 4.51.0 on single-H100 instances. The release includes the calibration manifest, out-of-fold scores,
 493 thresholds, overlap exclusions, benchmark hashes, and scripts that rebuild the figures and metric
 494 summaries from judged rows.