
When Does a Speech Model Know to Hand Off? Reading Competence from Hidden States for Real-Time Escalation

Anonymous Author(s)

Affiliation

Address

email

Abstract

Real-time speech systems answer most queries with a local model, but some queries call for a stronger expert. Selective escalation is well studied for text; speech tightens the constraint: the decision must be made before the model commits to a response, and the signal must survive distribution and modality shift. We ask whether a speech model’s own hidden states already carry such a competence signal. Our key finding is that they do, but not where escalation methods usually look: the conventional final-layer read collapses under distribution shift, while a mid-network read stays reliable, transfers from text-calibrated probes to speech input, and can be taken once at the model’s own speak commit, before the answer is produced, in about 30 ms. A linear probe on this representation gates a native full-duplex escalation loop: it raises delivered answer accuracy from 41% to 63% at 52% escalation while beating matched-rate random escalation, and, with weights, features, and layer fixed, improves all five external speech-QA benchmarks over always-local operation. These results identify a mid-network representation, rather than the final-layer read, as the practical signal for deciding when a speech model should hand off.

1 Introduction

Real-time spoken interaction leaves only a few hundred milliseconds between turns [Stivers et al., 2009]. Systems that meet this deadline keep a local speech model as the talker, so that the audio context and the voice stay in one session, and call a stronger remote model on turns the local model cannot answer. In a text interface this is a standard routing problem. In a speech interface the decision must also arrive before the talker commits to a response; otherwise the system must delay the expert call or retract an answer already under way. The central question of this paper is whether a duplex-trained speech model exposes its own likely failure early enough to hand the turn to a stronger model.

Part of the answer exists on the text side. Correctness of a text language model can be estimated from verbalized self-assessment, output uncertainty, or hidden states [Kadavath et al., 2022, Azaria and Mitchell, 2023, Burns et al., 2023, Kossen et al., 2024], and text routers use query-side or pre-generation features to choose a model [Ong et al., 2024, Varshney et al., 2026, Lugoloobi et al., 2026]. On the speech side, duplex models learn interaction control as a generation policy: special or synchronization tokens decide when the model should listen or speak [Wang et al., 2024a, Zhang et al., 2024, Veluri et al., 2024, Fu et al., 2024, Yu et al., 2024], and recent systems extend the same




Figure 1: **The deployed system.** (1) Speech streams in 1 s chunks into a full-duplex local model that decides for itself when to speak. (2) At that speak commit, a mid-network (L22) linear read scores the turn in ~ 30 ms, before the answer is produced. (3) Lower-risk turns stay local; higher-risk turns are sent to a cloud expert while the talker stalls. (4) The returned answer is spoken verbatim in the talker’s own voice. Delivered accuracy on the internal test set rises from 40.8% to 62.9% at 52% escalation in native full-duplex sessions.

33 mechanism to hand work to a reasoning layer or a retrieval service [Huang et al., 2026, Chien et al.,
34 2026].

35 The trained-control path does not close the gap, for two reasons. First, its cost: the policy is trained
36 into the weights, which takes large amounts of duplex-specific data (hundreds of thousands of hours of
37 synthetic dialogue for SyncLLM, millions of hours of audio for Moshi [Veluri et al., 2024, Défossez
38 et al., 2024]) and risks eroding the backbone’s language ability, which is the stated motivation for
39 frozen-backbone designs [Wang et al., 2024b, Yu et al., 2024]. Second, the interface is wrong for
40 escalation: a speak token decides when the model may talk, not whether its answer will be right,
41 and it arrives as a binary action at the moment of commitment. Even when the token triggers a
42 reasoning step instead of speech [Huang et al., 2026], the decision is still a trained binary action: it
43 offers no calibrated score to threshold at a fixed expert budget, and the policy cannot change without
44 retraining. What remains unknown is whether the escalation decision needs training at all: whether a
45 frozen duplex-trained checkpoint already carries a competence signal that can be read before response
46 commitment. Text-probe results make this plausible, but they do not say where such a read stays
47 reliable after duplex speech training, whether a probe calibrated on inexpensive text data still reads
48 speech-input states, or whether the read arrives early enough to act on.

49 These three conditions are exactly what a deployment needs: a score that separates individual failures
50 rather than memorizing the query types seen in calibration, that survives the text-to-speech shift, and
51 that improves which turns are escalated at a fixed call budget. If such a signal exists, any frozen
52 duplex model gains selective escalation for the cost of a linear probe; if it does not, speech escalation
53 inherits the training costs above.

54 We test these conditions on MiniCPM-o 4.5 [Cui et al., 2026], a 9B omni model with full-duplex
55 training, against matched controls: raw backbones, streaming-omni and audio-only fine-tunes, a
56 frozen-backbone adapter, and a second duplex generation. The study uses judged failure labels
57 on 600 public-benchmark queries in five frozen pools, leave-one-pool-out (LOPO) transfer as the
58 primary robustness test, matched text and speech inputs, and a complete live escalation loop. We
59 compare decode-time uncertainty, verbalized self-assessment, and linear hidden-state probes across
60 layers, then freeze the winning read for evaluation on public speech benchmarks. Model checkpoints
61 stay frozen; only linear probes are fit on stored activations. Throughout, duplex-trained refers to
62 the weights; the primary live evaluation is turn-based, with audio input and text output, and the
63 concurrent and native full-duplex serving regimes are validated separately (Appendices H and O).

64 This study makes three contributions:

1. **Transfer moves the readable layer.** The conventional final-layer probe holds in distribution but inverts on held-out math (AUC 0.372); the same last-token read at a mid-network layer reaches 0.931. Matched comparisons tie the late-layer failure to the duplex checkpoints we test; its mechanism remains open (Appendix D).
2. **The mid-layer read satisfies the deployment conditions.** A text-calibrated probe at the same layer reads speech-input states, reproduces on real human recordings, and is available at the end of the user turn in ~ 30 ms, before the first output token.
3. **The read is sufficient to drive escalation.** In native full-duplex sessions it raises delivered accuracy from 40.8% to 62.9% at 52% escalation on the internal test set, beats matched-rate random escalation at the two upper tiers, and dominates the model’s own built-in reasoning mode on both accuracy and latency (Appendix F). With its layer and features frozen, the gate improves all five external speech-QA benchmarks over always-local operation, by 3 to 34 points at the aggressive tier, although selectivity over random is limited where the local floor is already high.

The claim is therefore about readout rather than mechanism: the model does not explain why it will fail, but its likely failure is more reliably readable mid-network than at the conventional output interface, early enough to support a live handoff.

2 Related Work

Pre-answer competence estimation. Text-language-model correctness can be estimated from verbal self-evaluation [Kadavath et al., 2022, Lin et al., 2022], output uncertainty, or hidden representations [Azaria and Mitchell, 2023, Burns et al., 2023, Marks and Tegmark, 2024]. Intermediate layers can carry more useful information than the output distribution [Orgad et al., 2025], and trained hidden-state probes are strong factual-confidence estimators [Mahaut et al., 2024]. Semantic entropy improves uncertainty estimation by comparing sampled answers [Kuhn et al., 2023, Farquhar et al., 2024]; Semantic Entropy Probes amortize this signal into a pre-generation or post-hoc hidden-state read [Kossen et al., 2024]. Attention-map probes and multimodal uncertainty fusion extend the same broad idea [Chuang et al., 2024, Zhang et al., 2026a]. This work establishes that competence is readable in text models. It does not show where the read remains reliable after duplex speech training or whether a probe calibrated on text transfers to speech-input states.

Selective routing. Text routing splits into two families by where the decision lives. In the first, the decision is read from outside the generator: HybridLLM [Ding et al., 2024], RouteLLM [Ong et al., 2024], and query-side routers [Xue et al., 2025, Hu et al., 2024, Li et al., 2026] choose a model for a complete text query from query features or preference data; FrugalGPT organizes cascades that can score an earlier model’s response [Chen et al., 2023], while AutoMix verifies a drafted answer before escalating [Aggarwal et al., 2024]. Closest to our signal, recent methods read pre-generation activations or fit linear success probes on frozen states [Varshney et al., 2026, Lugoloobi and Russell, 2025, Lugoloobi et al., 2026]. In the second family, the decision is trained into the generator: Toolformer and Self-RAG learn tool or reflection tokens [Schick et al., 2023, Asai et al., 2024]; Co-LLM and CITER learn token-level deferral [Shen et al., 2024, Zheng et al., 2025]; and pause, thought, or think/no-think tokens allocate more internal computation [Goyal et al., 2024, Zelikman et al., 2024, Yang et al., 2025, DeepSeek-AI et al., 2025]. Learning-to-defer and early-exit methods make related decisions within a model [Madras et al., 2018, Schuster et al., 2022]. Our contribution is not the linear pre-generation probe by itself. We test whether this kind of read survives duplex training and audio input, and use a deployment-time threshold without modifying the speech checkpoint.

Duplex control and live handoff. End-to-end spoken dialogue includes native duplex models [Nguyen et al., 2023, Défossez et al., 2024, Zeng et al., 2024, Fang et al., 2026], streaming omni models [Xu et al., 2025, Cui et al., 2026], frozen-backbone adapters [Wang et al., 2024b], and audio

fine-tunes [Chu et al., 2024]. Most control work in these systems concerns the floor: state, idle, synchronization, or interruption tokens [Wang et al., 2024a, Zhang et al., 2024, Veluri et al., 2024, Ma et al., 2024, Fu et al., 2024, Yu et al., 2024], structured action channels [Zhang et al., 2026b], or classifier heads over a frozen or lightly adapted backbone [Wang et al., 2024b, Chen et al., 2025, Liao et al., 2025]. Training this control is costly enough that reducing the cost is itself a research goal [Xu et al., 2024], and trained turn-taking remains unreliable across models [Lin et al., 2025]. These mechanisms decide when the model may speak; even the most recent native design adds only a few turn tokens to the vocabulary and answers every turn locally, whatever the model’s competence on that turn [Fang et al., 2026].

The direct neighbors of our gate are trained think triggers. DuplexOmni emits a think-control token at the end of a user utterance to hand work to a separate reasoning layer [Huang et al., 2026]; chronological thinking interleaves reasoning with the speech stream [Wu et al., 2026]; MiniCPM-o 4.5 ships a built-in reasoning mode [Cui et al., 2026]; MoshiRAG triggers asynchronous retrieval from the model’s emerging output [Chien et al., 2026]. These systems show that a live handoff is possible. But the trigger is trained into the weights from purpose-built duplex data, so changing the policy or the downstream resource means retraining; it is a binary action with no calibrated score, so the escalation rate cannot be set to a call budget; and the trained policies are hard to reuse or compare, because released artifacts often include the training recipe but not the trained checkpoint [Huang et al., 2026]. None of these works evaluates a competence score read from a frozen checkpoint before generation. Our probe replaces the trained trigger with such a read: it needs no training of the speech model, exposes a thresholdable score, and on our target model it dominates the built-in reasoning mode on both accuracy and latency (Appendix F).

Positioning. Prior work therefore establishes three ingredients separately: competence can be read from text models, text queries can be routed, and duplex models can coordinate speaking or external computation. To our knowledge, the reviewed literature does not establish their conjunction: whether a frozen duplex-trained speech model exposes a pre-answer competence signal that survives unseen query pools and text-to-speech transfer, and whether that signal arrives early enough to improve live selective escalation. We evaluate these conditions directly. The claims are about readout robustness and system utility; the training mechanism behind the late-layer failure remains open.

3 Reading Competence from Hidden States

This section establishes where and when a competence signal can be read from the frozen checkpoint. The analyses use the target model and matched controls introduced above: 600 public-benchmark queries in five frozen pools with judged failure labels, and logistic probes fit on stored hidden states. Full details of models, data, labels, and probe fitting are in §4.

3.1 Signal comparison under distribution shift

Table 3 compares the candidate signals on MiniCPM-o 4.5. In distribution the final-layer probe looks healthy (AUC 0.822), but most of that accuracy reflects query type rather than instance difficulty: the same hidden state predicts which pool a query came from, a classifier given only the pool identity recovers most of the probe’s margin, and appending true pool labels to the probe’s features adds nothing. The same decomposition holds at scale on the deployed calibration mix (Appendix O). We therefore treat leave-one-pool-out (LOPO) transfer as the primary metric, and under it the picture reverses: with text input the final-layer read falls to AUC 0.372 on held-out math, and when the trap pool is held out the probe scores it on the confident side even though the model fails every query in it. A router built on query text shows the same dependence on pool structure and also degrades under LOPO, with little improvement from substantially more data (Appendices B, E). In-distribution routing accuracy alone therefore cannot establish instance-level transfer.

Verbalized self-assessment transfers differently. Asked before answering, $p(\text{True})$ catches exactly the trap pool the probe misses (0.945). Asked after drafting, the model endorses its own confident-wrong

Table 1: Main results: selective escalation on the internal test set and five public speech-QA benchmarks (delivered answer accuracy, %, native full-duplex sessions, one live arm per tier). The probe is fixed; tier thresholds are label-free calibration quantiles per language (the internal tiers realize 8/25/52% escalation; realized rates per pool are in Appendix J). Matched-rate random mixes the always-local and measured always-escalate arms at each pool’s realized rate. avg and Δ (average gain over always-local) cover five external pools in the live block and four in the NVDA block. The relay-free bound scores the expert’s returned text directly. The bottom block evaluates our latest three-layer NemotronLabs-VoiceChat-11B probe by re-mixing cached local and expert outcomes after replay through its native frame protocol (Appendix C). AlpacaEval, reported separately on its 1–5 scale, changes from 3.68 to 4.01 at 44% escalation on MiniCPM (matched random 3.91; always-escalate 4.20 against 4.95 for the expert text, a spoken-length cap); the pass-3 NVDA ranking changes from 3.55 to 4.51 versus 4.23 for random.

	Internal	TriviaQA	WebQ	Llama Q.	SD-QA	Reason. zh	avg	Δ
<i>native full-duplex live sessions (delivered accuracy)</i>								
always-local	40.8	62.8	52.8	82.4	48.0	51.0	59.4	—
+ gate, conservative	46.7	60.4	59.2	78.8	54.0	62.4	63.0	+3.6
+ gate, balanced	54.2	70.0	62.4	82.4	65.0	66.3	69.2	+9.8
+ gate, aggressive	62.9	88.4	75.2	85.6	82.5	72.8	80.9	+21.5
matched random, aggressive	56.1	81.5	68.9	84.1	73.7	67.5	75.2	+15.8
always-escalate (measured)	70.4	96.0	79.6	91.6	87.0	83.7	87.6	+28.2
expert text as returned (relay-free bound)	76.5	96.4	82.3	93.2	88.5	86.1	89.3	+29.9
<i>NVDA three-layer answer-onset probe (pass-3 native-frame replay)</i>								
always-local	—	35.8	37.6	70.8	31.0	—	43.8	—
+ gate @15%	—	48.4	48.8	80.7	42.0	—	55.0	+11.2
+ gate @30%	—	61.0	59.2	85.8	54.0	—	65.0	+21.2
+ gate @50%	—	76.8	69.6	89.7	66.5	—	75.7	+31.9
matched random @50%	—	65.7	60.2	82.0	60.0	—	67.0	+23.2
always-expert (ceiling)	—	95.5	82.8	93.1	89.0	—	90.1	+46.3

161 answers and the signal inverts: post-draft $p(\text{True})$ is the strongest single readout in distribution but
 162 no longer provides a pre-answer routing signal.

163 3.2 Layer and position sweep

164 Matched comparisons tie the failure to the duplex checkpoints rather than to the backbone alone. The
 165 raw Qwen backbone, subsampled to match MiniCPM-o 4.5’s per-pool failure rates, holds LOPO math
 166 near 0.96 where the duplex checkpoint scores 0.372, and the Qwen2.5 family shows the same ordering,
 167 raw above streaming-omni above duplex (Table 4, appendix). The Qwen2-Audio comparison is
 168 not attribution evidence because the audio fine-tune also changes task capability and the failure
 169 distribution substantially (Appendix B). A layer $\times$ position sweep localizes the difference (Figure 2):
 170 it is concentrated at the last-token position in late layers on text input, mean-pooled reads degrade far
 171 less, and at L22 the last-token probe reaches **0.931** on the same held-out math pool, close to the raw
 172 backbone.

173 The failure is input-specific rather than universal: with audio calibration and audio input, the final-
 174 layer probe reaches 0.936 on LOPO math. We use L22 because one read point must remain robust
 175 for shifted text queries and support text-to-speech transfer, not because the final-layer signal is absent
 176 on the model’s native audio input.

177 The sweep establishes where the transferable read survives, not why it weakens late. Direct tests do
 178 not support the simple hypothesis that turn-control state takes over the last-token representation: the
 179 layer at which performance drops does not move in speak mode, and the inverting probe direction
 180 aligns with neither the control vocabulary nor an explicit listen/speak direction. A decomposition
 181 further shows that a distributed transferable component weakens over the final layers of the two
 182 duplex checkpoints but not their raw backbones. These observations constrain the explanation but do
 183 not identify a training objective or causal mechanism; the full audit is in Appendix D.

184 The mid-network pattern also appears outside the MiniCPM family. On NemotronLabs-VoiceChat-
 185 11B, a hybrid-Mamba duplex model, our latest three-layer probe reads L26/L30/L34 and reaches




Figure 2: The core representational result: LOPO math AUC by layer and position. On the duplex model the standard (final-layer, last-token) read inverts on text input while mid-network layers hold ≥ 0.9 ; raw backbones show no cliff. The deployed gate reads L22.

0.818 OOF AUC and 0.816 mean external AUC under pool-aligned judges in replay through the model’s native frame protocol (Appendix C).

3.3 Modality transfer

Probes trained on text hidden states and scored on speech-input hidden states show where the representation becomes shared across modalities: early layers are modality-specific, and from about half depth the transfer plateaus (0.855 at the deployed L22). This replicates on MiniCPM-o 2.6 and cross-family on Qwen2.5-Omni. The decisive control is Freeze-Omni: identical frozen LLM weights with audio attached by an adapter, and cross-modal transfer stays far below the end-to-end models. Text calibration therefore transfers in the end-to-end speech models tested here, but not in the adapter-attached control. Verbalized introspection shows the complementary failure: on audio input, pre-answer $p(\text{True})$ collapses on both MiniCPM duplex generations while both non-duplex controls stay intact. Controlled arms show the verbal self-check depends on the text-token pathway (duplicating the question as text restores it), while a linear probe remains usable on the same forward pass, and the restored self-check is additive with the probe (Appendix B, Table 7). The mid-layer band also replicates on 200 real human recordings from SD-QA, so the result is not limited to the TTS rendering. The live system keeps L22 and adds two pooled positions to form the deployed probe (Appendix B).

3.4 Decision timing and escalation budget

Scoring latency. The gate is a linear read on L22 states the forward pass already computes: a 4096-d dot product per position (the deployed probe adds two pooled positions at no added scoring latency; Appendix B). On an H100, measured from prefill start with CUDA synchronization, the score is available at P50 in 20 ms for text and 45 ms for audio, against a time to first token of 36/68 ms on the same benchmark. The measurement excludes network time, but it establishes the ordering the system needs: the cloud request can start before local decoding commits to an answer. Waiting for decoded tokens does not improve the read in our data and postpones the expert request (Appendix F).

Read position in the audio loop. In the streaming loop, scores taken from intermediate audio chunks are weaker than a read at the end of the turn and can miss information in the last chunk. The system therefore reads once, at the end of the user turn: after the final audio chunk it prefills the assistant turn and reads L22 at its start. In a 360-query replay, this single read reaches AUC

0.887 in about 30 ms, against 0.843 for the offline reference, and arrives before the first output token (Appendix G).

From score to escalation budget. Thresholding the score converts it into a call budget. On the frozen test split, the probe beats matched-rate random escalation at every operating point and recovers over half of the local-to-expert gap with a third of the expert calls (Figure 7, appendix). In distribution it is statistically tied with a pool-identity classifier, as §3.1 predicts; its advantage appears under transfer, where pool identity stops being sufficient. The gate is fit to local failure, not expert benefit: a benefit label would specialize the ranking to the current expert. The frozen scores remain useful under the stricter benefit label, but pools with little expert headroom remain a system-level limit (Appendix I). Comparisons with the model’s built-in reasoning mode are in Appendix F.

4 Experimental Setup

4.1 Models

The target model is MiniCPM-o 4.5 (9B; omni + full-duplex fine-tune of Qwen3-8B), run in bf16 with greedy decoding and seed 42. To attribute effects to the kind of fine-tune rather than the backbone, we assemble ten models spanning raw backbones, a second duplex generation (MiniCPM-o 2.6), a streaming-omni fine-tune (Qwen2.5-Omni-7B), an audio-only fine-tune (Qwen2-Audio-7B), and a frozen-backbone adapter model (Freeze-Omni, LLM weights byte-identical to its raw control); the full roster is Table 2 (appendix). The comparison statistic is always the within-pair difference. All model checkpoints remain frozen throughout; the escalation expert is gpt-5.5 at low reasoning effort.

4.2 Dataset and labels

The internal dataset is 600 queries in five pools, frozen before any gate experiment (360 calibration / 240 test): *trap* = SimpleQA [Wei et al., 2024] (the target fails 100%), *hard-knowledge* = MMLU-Pro [Wang et al., 2024c], *easy-fact* = TriviaQA [Joshi et al., 2017], *easy-chat* = Dolly/Alpaca [Conover et al., 2023, Taori et al., 2023], *hard-math* = GSM8K tail + MATH-500 [Cobbe et al., 2021, Hendrycks et al., 2021, Lightman et al., 2024]; public benchmarks only, no LLM-generated queries. Adequacy labels come from an LLM judge [Zheng et al., 2023] with structured outputs, blind to the gate and validated by a stronger re-judge (agreement 96.2% text / 94.5% audio; every headline number moves ≤ 1 point under the stronger labels). The audio arm renders the same frozen pool to speech with single-voice TTS (content unchanged; only modality varies), following Spoken-SQuAD / VoiceBench practice [Li et al., 2018, Chen et al., 2024]. Judge prompts, pins, and spend: Appendix M–N.

4.3 Probe training and baseline signals

We compare three signal families: (i) decode scalars: max token entropy or margin over the first k decode steps; (ii) linear probes: logistic regression on hidden states, with layer and position selected on calibration data; and (iii) verbalized self-evaluation [Kadavath et al., 2022]: $p(\text{True})$ -pre asks “would you answer this correctly?” before answering, $p(\text{True})$ -post asks about a drafted answer; $P(\text{Yes})$ is read from the first token’s distribution. The representation analyses use a single-position, 4096-d L22 probe fit on the 360-query calibration split. The live system uses the same layer plus two pooled positions, fit on a wider 2,310-query calibration mix sourced separately from the external benchmarks (a duplicate audit found 8 residual question collisions across the five pools; removing them changes no result — Appendix O). Concurrent and native serving change the read point, so those validations refit the same feature set in regime; Appendix B records the full calibration ledger.

4.4 Benchmarks and evaluation protocol

Audio-side findings are validated on real human recordings (SD-QA, Faisal et al., 2021); the frozen gate is evaluated on external public speech benchmarks with their official audio (§5.2). Following




Figure 3: Live escalation on the internal test set ($n=240$), native full-duplex sessions. Blue: delivered accuracy at each tier’s realized escalation rate (filled points are separate live arms; bars are bootstrap 95% CIs). Dashed grey: the matched-rate random reference between the never and always arms. Green dotted: the expert’s returned text scored directly, the relay-free bound. Selection beats random by 6.0 and 6.8 points at the balanced and aggressive tiers.

§3.1, we use LOPO AUC to test distribution transfer and text-trained/audio-scored AUC to test modality transfer. For the live system, the key comparison is against matched-rate random escalation: an improvement over always-local operation can come from the expert alone, while an improvement over random at the same call rate measures the value of the gate’s ranking.

4.5 Live loop and escalation path

The live loop is the model’s native full-duplex mode under its official serving configuration: audio streams in 1 s chunks into the generating context, and the listen/speak head decides for itself when to answer. The gate reads the three-position probe of §4.3 once, at that speak commit, on the 5,228-row calibration merge refit in regime with labels judged on the model’s own answers (frozen-test AUC 0.876). Escalation runs in three tiers from conservative to aggressive; each tier’s threshold is a quantile of the probe’s calibration score, set per language, so a tier sets the call budget without correctness labels; a second linear head on the same read exempts floor-management turns (“stop”, backchannels). If a turn escalates, a canned stall phrase in the talker’s own voice plays, the raw question audio is transcribed and sent to gpt-5.5, and the returned answer is spoken verbatim through the talker’s teacher-forced voice; the chunk loop keeps running, so the relay stays interruptible. Every session in §5.1 is one fresh duplex session with real transcription, expert, and relay, paced in wall time. The earlier turn-based harness loop, injection controls, and prefetch experiments are in Appendix G.

5 Results

5.1 Live escalation

We ran the 240 frozen test queries and the five external pools through five live arms each — never, three gate tiers, always — every arm a separate native full-duplex session per query

(§4.5). We score delivered answer accuracy: the correctness of what the user hears after the full path, judged by each benchmark’s own judge. On the internal set accuracy rises monotonically, 40.8% → 46.7% → 54.2% → 62.9% at 0/8/25/52% realized escalation, and the balanced and aggressive gains are significant (paired bootstrap; Table 1, Figure 3). The gate exceeds matched-rate random escalation by 6.0 and 6.8 points at those two tiers; the conservative tier, firing on 8% of turns, is within noise of random.

A measured always-escalate arm reaches 70.4% against 76.5% for the expert’s returned text scored directly, so the spoken relay now costs 6 points on this pool (a 400-character spoken cap on long answers); on the five external pools the two differ by 1 to 3 points. Local decoding is stochastic, and a zero-escalation re-run moves a pool’s floor by up to 3.6 points, so no claim rests on a smaller difference. Full tables, latency profiles, and boundary cases are in Appendix G.

Serving-regime checks. Replaying the loop with the talker enabled leaves the frozen ranking nearly unchanged (AUC 0.871/0.856 versus 0.866 with the talker disabled). A teacher-forced concurrent-prefill variant is less stable and requires an in-regime refit; its external results remain exploratory (Appendix H). Under the model’s native full-duplex head, the read moves to the head’s own listen→speak commit point, and the probe is refit in regime on a widened 5,228-row calibration merge under the official serving configuration, with labels judged on the native answers. The refit reads at AUC 0.876 internally and 0.771 on the four English external QA pools. The Mandarin reasoning pool is lower, at 0.621; a 1,500-question held-out Mandarin set of facts, knowledge, and commonsense reads at 0.811, so the gap is multi-step reasoning rather than language (Appendix O). Re-mixing the native gate’s decisions with cached expert outcomes beats matched-rate random on all five English pools at both deployed tiers; with per-language thresholds the Mandarin pool escalates at the nominal rates and beats matched-rate random at the two lower tiers. These checks support the ranking beyond the primary turn-based loop, but they also show that read points, probe weights, and thresholds must be calibrated for the serving regime (Appendix O).

The relay, not the gate, was the first bottleneck. The first native pass used the talker to voice the returned answer from a steering prompt. On Speech TriviaQA the expert’s returned text scored 0.960 under the official judge, but what the talker voiced after a steering prompt scored 0.728 (never-escalate 0.628), because the head truncated, ignored, or re-answered the relayed text on about a quarter of turns. Speaking the cleaned expert text verbatim through the talker’s own teacher-forced voice, with a context note in place of the steering prompt, closes that gap (same 60 queries: 0.733 versus 0.933) at about two seconds of added synthesis time, and is now the deployed relay, and all results in Table 1 use it. Native local floors sit 1–8 points below the earlier turn-based harness loop (Appendix G), but the routing gain is larger: at the aggressive tier delivered accuracy rises by 26 points on Speech TriviaQA, 22 on Speech WebQ, 34 on SD-QA, and 22 on the Mandarin reasoning pool, which now escalates at its nominal rates under per-language thresholds; Llama Questions gains 3 points at 19% escalation. Open-ended generation remains the boundary: on AlpacaEval the always-escalate arm scores 4.20 against 4.95 for the expert text, because the spoken cap removes half of the long answers. Classifying the 600 native failures by cause shows where the pre-answer read works: specific-but-wrong answers are 70% of failures and are ranked at AUC 0.81 (79% escalated at the aggressive tier against a 31% false-fire rate on correct turns), misheard questions at 0.89, whereas execution slips in multi-step reasoning are at chance (0.47) — the error has not happened yet when the gate reads, so only a later read point could recover them.

When the signal arrives. To measure that directly we re-dumped the full calibration merge and the six pools under the serving configuration and read the same layer-22 vector at five moments of one generation: one step before the commit, at the commit (deployed), and after the first, second, and third answer chunks, fitting one probe per moment with the deployed recipe. Across the five external pools the mean AUC is 0.753 before the commit, 0.755 at it, 0.762, 0.791, and 0.801 after one, two, and three answer chunks; the Mandarin reasoning pool moves from 0.648 to 0.769 after two chunks while the internal test set stays at 0.85. The commit-point state carries retrieval confidence; execution confidence only exists once the answer is under way. The pre-commit ceiling we report

is therefore a ceiling on the former, and the roughly 0.05 above it is available at the price of two to three seconds of a partly voiced wrong answer, a product trade rather than a probe improvement. Reading one step earlier, before the first spoken unit enters the state, helps Mandarin (+0.06) and costs Speech WebQ (−0.06): the first unit is often the answer entity in English and a discourse marker in Mandarin. A two-stage variant that decides at the commit for confident turns and defers only a gray band to the later read recovers under one point of mean delivered accuracy at a fixed budget (0.711 to 0.718 at 30%), with the gain confined to the Mandarin reasoning pool (+4.5 points), so the single commit-point decision remains the deployed form.

5.2 Transfer to public speech benchmarks

The shortcut analysis of §3.1 makes external transfer necessary. We freeze the probe weights, feature set, and layer selected on the internal calibration mix, then run the live loop of §4.5 on six public speech benchmarks with audio input and text output: Speech TriviaQA, Speech Web Questions, Llama Questions, and Reasoning QA from OpenAudioBench [Li et al., 2025] (spoken renderings of Joshi et al., 2017, Berant et al., 2013, and the set of Nachmani et al., 2024), SD-QA (real dialectal human speech) [Faisal et al., 2021], and VoiceBench AlpacaEval [Chen et al., 2024, Taori et al., 2023]. All audio is the benchmarks’ official pre-rendered audio; no evaluation query is used to fit the probe. The only per-pool adjustment is label-free: tier thresholds are quantiles of that pool’s score distribution, requiring no correctness labels or judge calls. This controls the call budget but cannot improve ranking. A single absolute threshold also beats matched-rate random on four of five QA pools, but its escalation rate varies from 2% to 48% (Appendix I). We use each benchmark’s official judge when available and verify that the local-only controls are on the published scale.

The frozen scores remain discriminative outside the calibration mix: AUC ranges from 0.683 to 0.806 across the five QA pools, with a mean of 0.771 (Appendix J). The gate also improves deployed-channel accuracy over always-local operation on every QA benchmark and at every budget; at 50% escalation, mean accuracy rises from 63.4% to 82.1%. These gains establish system utility but do not by themselves establish selectivity, because random escalation also adds expert answers.

The NVDA probe reaches 0.818 OOF AUC and 0.816 mean external AUC under each pool’s assigned judge. It averages linear heads at L26/L30/L34 on a pass-3 answer-onset capture from the model’s native frame protocol. This tests method transfer without running a live loop. The full training protocol and read-point ablations are in Appendix C.

Llama Questions provides the cleanest fixed-budget test. It contains a single query type, the kept-local subset scores 97.6% while the escalated subset scores 69.6% ($z=6.46$), and expert answers improve the latter to 88.8% ($p<.0001$). At 50% escalation, the measured system reaches 94.8%, compared with 93.6% for always escalating. Matched-rate precision reaches 95% of its base-rate cap (Appendix J). Cross-lingual transfer also improves when the English calibration mix is widened: Chinese Reasoning QA rises from AUC 0.621 to 0.683 despite adding no Chinese calibration data.

The transfer results also expose the boundary between predicting local failure and creating expert benefit. SD-QA is a positive real-speech case and exceeds matched-rate random at all three budgets. Web Questions is less stable, with little gain over random at the low and high budgets; in its top-scoring 15%, many local failures are also missed by the expert (Appendix I). AlpacaEval is a clearer negative: its open-ended failures do not present the retrieval-failure pattern the probe reads. On NVDA, the pass-3 probe beats matched-rate random on all four binary QA pools at 30% and 50% escalation ($p \leq .0008$); at 15%, SD-QA is inconclusive ($p=0.063$). On AlpacaEval, the same pass-3 ranking reaches 4.51 versus 4.23 for random at 50% ($p < .001$). Thus the three MiniCPM negatives all separate at the 30% and 50% budgets on the second family: external utility depends on both the competence ranking and whether the selected failures are repairable by the expert (Appendices C, I, and J).

6 Discussion

Interpretation. The competence signal this system runs on was already in the frozen weights. Nothing was trained into the speech model, and the read costs one dot product on states the forward pass computes anyway, so a deployed duplex model can gain selective escalation without new duplex data, a new control token, or a retrained policy. What matters is where the read is taken. The final layer is the default choice because it is the interface to the output, and on these duplex checkpoints it is the wrong one: the same linear read that inverts at the output holds up in the middle of the network. This also points back at duplex training itself. The matched comparisons show that the fine-tune, not the backbone, is what breaks the late-layer read, so a training recipe that adds interaction control may be degrading the model’s own competence interface while leaving its answers intact. We do not know the mechanism, and a model that reads its own state before answering still cannot explain why it will fail.

Scope of the signal. The probe is strongest on retrieval-like failures, where the state reflects a missing answer before generation. It is weak on false premises and complementary to signals that appear during decoding (Appendix G). This boundary explains why failure AUC and system benefit are not the same quantity. A correct ranking cannot help when the expert fails on the same turns, and an open-ended quality benchmark may not contain the factual-retrieval event the probe detects. The unstable Web Questions curve and the negative AlpacaEval result are part of that boundary rather than exceptions to it; SD-QA, by contrast, confirms selective gains on real human speech. Selective escalation needs both a readable local failure and expert headroom.

Calibration. The model weights stay frozen, but the gate is still a fitted component. Text-only calibration is enough for the cross-modal result on end-to-end speech models, the deployed probe uses a wider benchmark-disjoint mix, and both the concurrent and native read points require refitting in their own regime. The method avoids fine-tuning the talker and moves the operating budget through a threshold, but it does not remove the need for representative calibration data.

Limitations. The internal study uses one phrasing per $p(\text{True})$ variant, and its real-speech validation is scripted read speech in a single dialect. Formula-symbolic content is out of scope for the audio arm, since any speech rendering destroys it. Labels come from one judge family, with no human labels; a stronger re-judge bounds the resulting shifts at about one point. Each query is run once per tier, which puts a replication floor of two to three points on every live accuracy, and no claim here rests on a smaller difference. The second-family NVDA result covers method transfer through the model’s native frame protocol, not matched-pair attribution. Its external pools were inspected during the variant review and thus need fresh-pool confirmation; its layer and features were selected on calibration-only LOPO metrics. For the main MiniCPM probe, the layer and features were chosen by cross-validation on the internal calibration split and then confirmed on held-out pools, rather than by nested selection inside each held-out pool. The mechanism behind the late-layer decay on text input remains unresolved (Appendix D).

Serving regime is the sharpest practical limit. Calibration does not carry across regimes: under concurrent prefill the frozen read degrades and needs an in-regime refit, and the native read needs its own refit, wider calibration, and labels judged on native answers before it matches the turn-based loop (Appendices H, O). That native read also appears saturated, in that more labels of the same kind, larger heads, and a distilled second-signal student all leave it close to where it is. Finally, the evaluation is scripted and single-turn. Disfluent correction, chained tool calls, and multi-turn grounding remain open [Lin et al., 2026], and in native testing a new question asked during the expert wait usually derails the pending relay, so a production system has to cancel or replace an in-flight request when the user moves on. Interruption inherits the duplex head’s own turn-taking, and here the harness’s bookkeeping matters more than decode-time levers. Our first design narrated escalations to the head in bracketed system notes and let the condemned local answer run on silently to its own end-of-turn; after a barge the head then committed to the following question only about two times in three, and no decode-side intervention (down-weighting the listen token, forced listen

windows, closing the interrupted turn) made the commit reliable. The failures traced to the context those mechanisms left behind—content the head never saw in training, stacked against speaking. Replacing them with teacher-forcing—whatever audio the system actually played (stall line, relayed answer up to the cut) is written into the context as the head’s own speech and closed with a turn-end token, the exact shape a native yield leaves—restored post-interruption commits in our live runs (8/8, vs. the prior $\sim 2/3$ ceiling). The residual is the head’s baseline turn-taking variance (a ~ 1 -in-6 chance of ignoring a first utterance in a clean session), which escalation neither causes nor can fix without an external voice-activity trigger of exactly the kind the native regime is meant to retire.

References

- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2310.12963.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.11511.
- Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. arXiv:2304.13734.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2013. WebQuestions.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2212.03827.
- Chi Seng Cheang, Hou Pong Chan, Wenxuan Zhang, and Yang Deng. Do LLMs really know what they don’t know? internal states mainly reflect knowledge recall rather than truthfulness. *arXiv preprint arXiv:2510.09033*, 2025.
- Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- Qian Chen, Yafeng Chen, Yanni Chen, et al. MinMo: A multimodal large language model for seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024.
- Chung-Ming Chien, Manu Orsini, Eugene Kharitonov, Neil Zeghidour, Karen Livescu, and Alexandre Défossez. MoshiRAG: Asynchronous knowledge retrieval for full-duplex speech language models. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2604.12928.
- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, et al. Qwen2-Audio technical report. *arXiv preprint arXiv:2407.10759*, 2024.
- Yung-Sung Chuang, Linlu Qiu, Cheng-Yu Hsieh, Ranjay Krishna, Yoon Kim, and James Glass. Lookback lens: Detecting and mitigating contextual hallucinations in large language models using only attention maps. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. GSM8K.

473 Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick
474 Wendell, Matei Zaharia, and Reynold Xin. Free dolly: Introducing the world’s first truly open
475 instruction-tuned LLM. Databricks blog, 2023. databricks-dolly-15k.

476 Junbo Cui, Bokai Xu, Chongyi Wang, Tianyu Yu, Weiyue Sun, Yingjing Xu, et al. MiniCPM-o 4.5:
477 Towards real-time full-duplex omni-modal interaction. *arXiv preprint arXiv:2604.27393*, 2026.

478 Joey David. Temporal predictors of outcome in reasoning language models. *arXiv preprint*
479 *arXiv:2511.14773*, 2025.

480 DeepSeek-AI, Daya Guo, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via
481 reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

482 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
483 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue.
484 *arXiv preprint arXiv:2410.00037*, 2024.

485 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Rühle, Laks
486 V. S. Lakshmanan, and Ahmed Hassan Awadallah. Hybrid LLM: Cost-efficient and quality-
487 aware query routing. In *International Conference on Learning Representations (ICLR)*, 2024.
488 arXiv:2404.14618.

489 Fahim Faisal, Sharlina Keshava, Md Mahfuz ibn Alam, and Antonios Anastasopoulos. SD-QA:
490 Spoken dialectal question answering for the real world. In *Findings of the Association for*
491 *Computational Linguistics: EMNLP 2021*, 2021. arXiv:2109.12072.

492 Qingkai Fang, Shoutao Guo, and Yang Feng. BayLing-Duplex: Native full-duplex speech dialogue
493 with a single autoregressive LLM. *arXiv preprint arXiv:2606.14528*, 2026.

494 Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large
495 language models using semantic entropy. *Nature*, 630:625–630, 2024.

496 Chaoyou Fu, Haojia Lin, Zuwei Long, Yunhang Shen, Yuhang Dai, Meng Zhao, et al. VITA: Towards
497 open-source interactive omni multimodal LLM. *arXiv preprint arXiv:2408.05211*, 2024.

498 Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh
499 Nagarajan. Think before you speak: Training language models with pause tokens. In *ICLR*, 2024.

500 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
501 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
502 *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks*, 2021.
503 arXiv:2103.03874.

504 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
505 Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing
506 systems. *arXiv preprint arXiv:2403.12031*, 2024.

507 Muye Huang, Lingling Zhang, Xingyu Yu, Lei Shi, Zhanyu Ma, Jun Xu, Jiuchong Gao, Jinghua Hao,
508 Renqing He, and Jun Liu. DuplexOmni: Real-time listening, seeing, thinking, and speaking for
509 full-duplex interaction. *arXiv preprint arXiv:2606.09186*, 2026.

510 Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly
511 supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual*
512 *Meeting of the Association for Computational Linguistics (ACL)*, 2017.

513 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas
514 Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly)
515 know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

516 Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal.
517 Semantic entropy probes: Robust and cheap hallucination detection in LLMs. *arXiv preprint*
518 *arXiv:2406.15927*, 2024.

519 Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for
520 uncertainty estimation in natural language generation. In *International Conference on Learning*
521 *Representations (ICLR)*, 2023. *arXiv:2302.09664*.

522 Chia-Hsuan Li, Szu-Lin Wu, Chi-Liang Liu, and Hung-yi Lee. Spoken SQuAD: A study of miti-
523 gating the impact of speech recognition errors on listening comprehension. In *Interspeech*, 2018.
524 *arXiv:1804.00320*.

525 Hao Li, Yiqun Zhang, Zhaoyan Guo, Chenxu Wang, Shengji Tang, Qiaosheng Zhang, Yang Chen,
526 Biqing Qi, Peng Ye, Lei Bai, Zhen Wang, and Shuyue Hu. LLMRouterBench: A massive
527 benchmark and unified framework for LLM routing. *arXiv preprint arXiv:2601.07206*, 2026.

528 Tianpeng Li, Jun Liu, Tao Zhang, Yuanbo Fang, Da Pan, Mingrui Wang, et al. Baichuan-Audio:
529 A unified framework for end-to-end speech interaction. *arXiv preprint arXiv:2502.17239*, 2025.
530 Introduces OpenAudioBench.

531 Borui Liao, Yulong Xu, Jiao Ou, Kaiyuan Yang, Weihua Jian, Pengfei Wan, and Di Zhang. FlexDuo:
532 A pluggable system for enabling full-duplex capabilities in speech dialogue systems. *arXiv preprint*
533 *arXiv:2502.13472*, 2025.

534 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
535 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
536 *Conference on Learning Representations (ICLR)*, 2024. *arXiv:2305.20050*; source of the MATH-
537 500 subset.

538 Guan-Ting Lin, Jiachen Lian, Tingle Li, Qirui Wang, Gopala Anumanchipalli, Alexander H. Liu, and
539 Hung-yi Lee. Full-Duplex-Bench: A benchmark to evaluate full-duplex spoken dialogue models
540 on turn-taking capabilities. *arXiv preprint arXiv:2503.04721*, 2025.

541 Guan-Ting Lin, Chen Chen, Zhehuai Chen, and Hung-yi Lee. Full-Duplex-Bench-v3: Benchmarking
542 tool use for full-duplex voice agents under real-world disfluency. *arXiv preprint arXiv:2604.04847*,
543 2026.

544 Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in
545 words. *Transactions on Machine Learning Research*, 2022. *arXiv:2205.14334*.

546 William Lugoloobi and Chris Russell. LLMs encode how difficult problems are. *arXiv preprint*
547 *arXiv:2510.18147*, 2025.

548 William Lugoloobi, Thomas Foster, William Bankes, and Chris Russell. LLMs encode their failures:
549 Predicting success from pre-generation activations. *arXiv preprint arXiv:2602.09924*, 2026.

550 Ziyang Ma, Yakun Song, Chenpeng Du, Jian Cong, Zhuo Chen, Yuping Wang, Yuxuan Wang, and
551 Xie Chen. Language model can listen while speaking. *arXiv preprint arXiv:2408.02622*, 2024.

552 David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and
553 accuracy by learning to defer. In *Advances in Neural Information Processing Systems (NeurIPS)*,
554 2018.

555 Matéo Mahaut, Laura Aina, Paula Czarnowska, Momchil Hardalov, Thomas Müller, and Lluís
556 Màrquez. Factual confidence of LLMs: on reliability and robustness of current estimators. In
557 *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.

558 Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language
559 model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
560 *arXiv:2310.06824*.

561 Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh
562 Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question
563 answering and speech continuation using spectrogram-powered LLM. In *International Conference*
564 *on Learning Representations (ICLR)*, 2024. arXiv:2305.15255; Spectron, source of the Llama
565 Questions dataset.

566 Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden
567 Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux.
568 Generative spoken dialogue language modeling. *Transactions of the Association for Computational*
569 *Linguistics*, 2023. arXiv:2203.16502.

570 Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez,
571 M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data.
572 *arXiv preprint arXiv:2406.18665*, 2024.

573 Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Koteck, and Yonatan
574 Belinkov. LLMs know more than they show: On the intrinsic representation of LLM hallucinations.
575 In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.02707.

576 Geoff Pleiss, Tianyi Zhang, Ethan Elenberg, and Kilian Q. Weinberger. Identifying mislabeled data
577 using the area under the margin ranking. In *NeurIPS*, 2020.

578 Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer,
579 Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use
580 tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2302.04761.

581 Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and
582 Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information*
583 *Processing Systems (NeurIPS)*, 2022. arXiv:2207.07061.

584 Shannon Zejiang Shen, Hunter Lang, Bailin Wang, Yoon Kim, and David Sontag. Learning to decode
585 collaboratively with multiple language models. In *Proceedings of the 62nd Annual Meeting of the*
586 *Association for Computational Linguistics (ACL)*, 2024. arXiv:2403.03870; Co-LLM.

587 Tanya Stivers, N. J. Enfield, Penelope Brown, Christina Englert, Makoto Hayashi, Trine Heinemann,
588 Gertie Hoymann, Federico Rossano, Jan Peter de Ruiter, Kyung-Eun Yoon, and Stephen C.
589 Levinson. Universals and cultural variation in turn-taking in conversation. *Proceedings of the*
590 *National Academy of Sciences*, 106(26):10587–10592, 2009.

591 Swabha Swayamdipta, Roy Schwartz, Nicholas Lourie, Yizhong Wang, Hannaneh Hajishirzi, Noah A.
592 Smith, and Yejin Choi. Dataset cartography: Mapping and diagnosing datasets with training
593 dynamics. In *EMNLP*, 2020.

594 Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy
595 Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following LLaMA model.
596 GitHub repository, 2023. zh split: shibing624/alpaca-zh.

597 Tanay Varshney, Annie Surla, Michelle Xu, Gomathy Venkata Krishnan, Maximilian Jeblick, David
598 Austin, Neal Vaidya, and Davide Onofrio. LLM router: Rethinking routing with prefix activations.
599 *arXiv preprint arXiv:2603.20895*, 2026.

600 Bandhav Veluri, Benjamin N. Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. Beyond
601 turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In *Proceedings of*
602 *EMNLP (Main)*, 2024. arXiv:2409.15594.

603 Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung,
604 Denny Zhou, Quoc Le, and Thang Luong. Freshllms: Refreshing large language models with
605 search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023.

606 Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. A full-duplex
607 speech dialogue scheme based on large language models. In *Advances in Neural Information*
608 *Processing Systems (NeurIPS)*, 2024a. arXiv:2405.19487.

609 Xiong Wang, Yangze Li, Chaoyou Fu, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni:
610 A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint*
611 *arXiv:2411.00774*, 2024b.

612 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, et al. MMLU-Pro: A more robust and
613 challenging multi-task language understanding benchmark. In *Advances in Neural Information*
614 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2024c. arXiv:2406.01574.

615 Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese,
616 John Schulman, and William Fedus. Measuring short-form factuality in large language models.
617 *arXiv preprint arXiv:2411.04368*, 2024. SimpleQA.

618 Donghang Wu, Haoyang Zhang, Chen Chen, Tianyu Zhang, Fei Tian, Xuerui Yang, Gang Yu, Hexin
619 Liu, Nana Hou, Yuchen Hu, and Eng Siong Chng. Chronological thinking in full-duplex spoken
620 dialogue language models. In *Proceedings of the 27th Annual Meeting of the Special Interest*
621 *Group on Discourse and Dialogue (SIGDIAL)*, 2026. arXiv:2510.05150.

622 Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, et al. Qwen2.5-
623 Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025.

624 Wang Xu, Shuo Wang, Weilin Zhao, Xu Han, Yukun Yan, Yudi Zhang, Zhe Tao, Zhiyuan Liu, and
625 Wanxiang Che. Enabling real-time conversations with minimal training costs. *arXiv preprint*
626 *arXiv:2409.11727*, 2024.

627 Leyang Xue, Yao Fu, Luo Mai, and Mahesh K. Marina. HybridServe: Efficient serving of large AI
628 models with confidence-based cascade routing. *arXiv preprint arXiv:2505.12566*, 2025.

629 An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

630 Wenyi Yu, Siyin Wang, Xiaoyu Yang, Xianzhao Chen, Xiaohai Tian, Jun Zhang, Guangzhi Sun,
631 Lu Lu, Yuxuan Wang, and Chao Zhang. SALMONN-omni: A codec-free LLM for full-duplex
632 speech understanding and generation. *arXiv preprint arXiv:2411.18138*, 2024.

633 Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, Nick Haber, and Noah D. Goodman.
634 Quiet-STaR: Language models can teach themselves to think before speaking. *arXiv preprint*
635 *arXiv:2403.09629*, 2024.

636 Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong,
637 and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot. *arXiv*
638 *preprint arXiv:2412.02612*, 2024.

639 Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno
640 Kriz. Unified multimodal uncertain inference. *arXiv preprint arXiv:2604.08701*, 2026a.

641 Haoyang Zhang, Jun Chen, Donghang Wu, Yuxin Li, Yuxin Zhang, Xiangyu Tony Zhang, Che Liu,
642 Qingjian Lin, Yizhou Peng, Hexin Liu, Eng Siong Chng, Chao Yan, Boyong Wu, Yechang Huang,
643 Xuerui Yang, and Fei Tian. DuplexSLA: A full-duplex spoken language model with synchronized
644 speech, language, and action. *arXiv preprint arXiv:2605.20755*, 2026b.

645 Xinrong Zhang, Yingfa Chen, Shengding Hu, Xu Han, Zihang Xu, Yuanwei Xu, Weilin Zhao,
646 Maosong Sun, and Zhiyuan Liu. Beyond the turn-based game: Enabling real-time conversations
647 with duplex models. In *Proceedings of EMNLP*, 2024. arXiv:2406.15718.

648 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang,
649 Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.
650 Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information*
651 *Processing Systems (NeurIPS), Datasets and Benchmarks*, 2023. arXiv:2306.05685.

652 Wenhao Zheng, Yixiao Chen, Weitong Zhang, Souvik Kundu, Yun Li, Zhengzhong Liu, Eric P.
653 Xing, Hongyi Wang, and Huaxiu Yao. CITER: Collaborative inference for efficient large language
654 model decoding with token-level routing. In *Conference on Language Modeling (COLM)*, 2025.
655 arXiv:2502.01976.

656 A Model roster

Table 2: Model roster. “Pair” groups a fine-tune with its raw backbone control; the comparison statistic is always the within-pair difference.

model	type	pair	role
MiniCPM-o 4.5	omni + duplex FT	Qwen3-8B	system target
Qwen3-8B	raw	—	pair control
MiniCPM-o 2.6	duplex FT	Qwen2.5-7B	duplex replication
Qwen2.5-Omni-7B	streaming omni FT	Qwen2.5-7B	omni control
Qwen2.5-7B	raw	—	pair control
Qwen2-Audio-7B	audio FT	Qwen2-7B	audio-FT control
Freeze-Omni	frozen LLM + adapter	Qwen2-7B	adapter control
Qwen2-7B	raw	—	pair control
Mistral-7B-v0.3	raw	—	backbone diversity
GLM4-9B-chat	raw	—	backbone diversity

657 B Extended signal and readout analyses

Table 3: Signals on MiniCPM-o 4.5, judged failure labels. In-dist = calibration-split AUC (probes 5-fold out-of-fold); LOPO math = leave-one-pool-out AUC on held-out math; audio = the same readout under speech input (final/mid-layer probe: audio-input LOPO math / text-only calibration read on audio states). Full per-signal detail: Appendix B.

signal	in-dist	LOPO math	audio input
max entropy@4	0.696	≈chance on raw backbones	—
$p(\text{True})$ -pre	0.807	trap 0.945 (probe: 0.232)	collapses
$p(\text{True})$ -post (full draft)	0.899	—	degrades
final-layer probe	0.822	0.372 (inverts)	0.936 (intact)
mid-layer (L22) probe	0.866	0.931	0.855 (text-calibrated)

658 **The in-distribution shortcut, quantified.** The numbers behind §3.1’s claim that the final-layer
659 probe’s in-distribution AUC of 0.822 is mostly pool identity: a linear read of the same hidden state
660 classifies a query’s source pool at 95.8% accuracy; a pool-identity oracle alone reaches AUC 0.715;
661 and appending true pool labels to the probe’s features moves it 0.822→0.821. In-distribution area
662 over random is statistically on par between the gate and the pool oracle (Appendix E); the two
663 separate only under LOPO, where the probe inverts to 0.372 and the oracle is undefined.

664 **ROC curves.** Figure 4 shows the calibration-split ROC for the signals of Table 3.

665 $p(\text{True})$ **backbone dependence.** Across backbones, post-draft $p(\text{True})$ beats the trained final-layer
666 probe on four of five raw/duplex models tested; the exception (GLM4-9B-chat, 0.685 vs. probe 0.798)
667 shows the pattern depends on the backbone’s self-evaluation quality.

668 **Matched pairs.** Table 4: with label coverage matched by subsampling the raw control to the duplex
669 model’s per-pool fail rates, the raw backbone keeps LOPO math at 0.825/0.809 while duplex models
670 sit at chance: label coverage is ruled out; the representation changed. Degradation is ordered raw
671 > omni-streaming > duplex. Table 5 gives the by-depth summary behind Figure 2: the within-pair
672 deficit at the best mid-layer is only −0.03 vs. −0.59 at the final layer; mean-pooled probes survive to
673 the final layer on duplex models (LOPO math ≈0.80 at L35 on o4.5, 0.674 at L27 on o2.6, 0.13–0.15
674 below each model’s mid-layer last-token peak, and uniformly weaker in-mix, −0.058 at L22); duplex
675 damage is severe but local, omni-streaming damage mild but diffuse.




Figure 4: ROC on the calibration split. $p(\text{True})$ -post leads but needs a full drafted answer; the probes and $p(\text{True})$ -pre are pre-answer.

Table 4: LOPO transfer (AUC, held-out pool) of the final-layer last-token probe, Qwen2.5-7B family.

LOPO pool	Qwen2.5-7B (raw)	Qwen2.5-Omni	MiniCPM-o 2.6	[o4.5 vs Qwen3]
hard-math	0.809	0.745	0.538 $\approx$ chance	0.372 (inverts)
hard-knowledge	0.680	0.613	0.526 $\approx$ chance	0.61

676 **Modality specificity of the cliff.** Under audio input the final layer does not invert (L35 audio LOPO
677 math 0.936 vs. 0.366 text); a speak-mode template control on text input leaves the cliff unchanged
678 (0.362): the operative variable is the modality of the input context, not the output mode.

679 **Audio collapse of verbalized introspection.** Table 6. Three controlled arms pin the mechanism
680 on MiniCPM-o 4.5: not perception (ASR audit WER 0.074; $p(\text{True})$ on the model’s own transcript
681 snaps back to 0.074; the well-heard subset still collapses); not a context prior (irrelevant filler audio
682 does not inflate p_{yes}); binding (duplicating the question as text alongside the audio fully restores
683 introspection, 0.034). The verbal instance-check runs over text-token pathways that audio embeddings
684 do not feed, while a linear probe reads the same information off the same forward pass at 0.93+.

Table 5: LOPO math transfer by depth (last-token probes).

model	best mid-layer	final layer	shape
Qwen3-8B (raw)	0.964 (L33/36)	0.958	plateau, no cliff
MiniCPM-o 4.5 (duplex)	0.931 (L22/36)	0.366	cliff, inverts
Qwen2.5-7B (raw)	0.893 (L21/28)	0.809	mild dip
Qwen2.5-Omni (streaming)	0.794 (L16/28)	0.746	diffuse depression
MiniCPM-o 2.6 (duplex)	0.822 (L21/28)	0.540	cliff

685 The failure implies its own cheap fix, confirmed twice: re-present the query as text before asking
686 (“repeat-then-judge”).

Table 6: Pre-answer introspection under audio input. Trap p_{yes} = mean stated probability of answering correctly on a 100%-fail pool (lower is better).

model	FT type	trap p_{yes} : text $\rightarrow$ audio	audio $p(\text{True})$ -pre AUC
MiniCPM-o 4.5	duplex	0.055 $\rightarrow$ 0.556 (collapse)	0.786 (trap dead)
MiniCPM-o 2.6	duplex	0.196 $\rightarrow$ 0.345	0.491 $\approx$ chance
Qwen2.5-Omni	omni-streaming	0.279 $\rightarrow$ 0.213 (intact)	0.727
Freeze-Omni	frozen + adapter	0.165 $\rightarrow$ 0.112 (intact)	0.612 (capability caveat)

687 **Probe $\oplus$ repeat-then-judge fusion.** The restored self-evaluation is additive with the probe. We
688 stack the deployed L22 probe with $p(\text{True})$ -pre asked over the model’s own transcript (one short text
689 prefill + a 1-token decode per turn; logistic stacker fit on the 360 calibration rows, native regime,
690 Table 7). On the internal test split the fusion is the first significant AUC lift over the deployed
691 probe (0.836 $\rightarrow$ 0.862; bootstrap 95% CI on Δ excludes zero), and it transfers to one external pool
692 (TriviaQA 0.789 $\rightarrow$ 0.820); the remaining English pools move positively but within CI, and the
693 Mandarin reasoning pool is untouched: there the self-evaluation itself carries no signal (AUC 0.586).
694 Contrary to our a-priori expectation, the gain tracks where the self-evaluation signal is strong, not
695 where the probe is weak; a diagnostic arm asking $p(\text{True})$ -pre on the original query text scores within
696 0.01–0.02 of the transcript arm on every pool, so transcription quality is not the limiter.

Table 7: Probe $\oplus$ repeat-then-judge $p(\text{True})$ -pre fusion (native regime, deployed read point). Δ = fusion – probe, bootstrap 95% CI.

pool	n	probe	rtj solo	fusion	Δ [95% CI]
internal test	240	0.836	0.836	0.862	+0.026 [+0.004, +0.048]
TriviaQA	250	0.789	0.761	0.820	+0.030 [+0.000, +0.062]
WebQ	250	0.745	0.723	0.771	+0.026 [−0.007, +0.059]
Llama-Q	250	0.686	0.674	0.691	+0.006 [−0.020, +0.031]
SD-QA	200	0.803	0.742	0.808	+0.004 [−0.023, +0.032]
Reasoning-zh	202	0.640	0.586	0.618	−0.021 [−0.076, +0.032]

697 **Real-speech validation.** The matched-pair design rerun on 200 real human recordings (VoiceBench
698 SD-QA, USA split) replicates all three audio-side findings: modality tax (0.400 $\rightarrow$ 0.450), audio
699 overconfidence (paired p_{yes} shift +0.089; failure subset 0.415 $\rightarrow$ 0.581), and the layer structure (text-
700 calibrated read on real-speech audio holds 0.76–0.80 at L22–L26, peak 0.800 at L25, transfer across
701 both content and modality).

702 **Synthesis.** End-to-end multimodal training builds the shared mid-layer core; duplex-style training
703 additionally damages the readouts; omni-streaming gets both right; an adapter alone gets neither
704 (probe readout text-fragile, verbal readout audio-fragile; in both cases knowledge survives and a
705 readout breaks, and in both cases the breakage is duplex-fine-tune-specific).

The deployed probe, and which calibration each result uses. The mechanistic probe is a single-position L22 read (4096-d), fit on the 360-query calibration split; every representational result (§3, Appendices B, D) is this probe. The deployed probe adds two zero-latency pooled positions (an 8-token tail mean and the running user-audio mean; 12,288-d), refit on a calibration mix widened to 2,310 public-benchmark queries sourced separately from every evaluation benchmark (overlap audit below); a pre-registered guard held (frozen internal test AUC 0.860→0.879) and feature selection never saw the externals. It drives the turn-based live loop (§5.1), the frozen transfer results (§5.2), and the talker-enabled replay, frozen throughout, calibrated on text and speech-rendered internal data only. The serving-regime validations are the exceptions. In the concurrent regime, the same 12,288-d feature set is refit in-regime on 360 rows for the loop of Table 14 and on the full 2,310-row mix to measure the effect of calibration scale (Appendix H). In the native regime, the same feature set is refit at the native listen→speak read point; the deployed native gate is the 5,228-row official-configuration merge of Table 17 (Appendix O). **Selection history.** L22 and the three-position feature set were chosen on internal calibration-split cross-validation before any external benchmark was scored (the 8-position × layer sweep and the 19-configuration feature sweep are both calibration-only), and the internal frozen test split was the pre-registered guard. The external pools were then read once with the frozen artifact. Two honest exposures remain: the layer band was picked on internal LOPO rather than by nested selection inside each held-out pool, and the second duplex family’s mid-band peak (L30–34 of 56) was located with its own sweep, so both layer choices are best read as confirmed on held-out pools, not as untouched-holdout estimates. In the mid-layer band the in-mix tradeoff optimum is broad (L20–L30) and the mid-layer probe’s in-distribution area advantage is modest: its real advantage is LOPO robustness, which an in-mix test cannot show.

728 C Second duplex family: NemotronLabs-VoiceChat-11B

Probe training protocol. Our latest probe uses the pass-3 replay through NVIDIA’s native frame protocol with the VoiceChat checkpoint frozen. Of 2,550 English calibration queries, 2,481 produce a commit frame and enter the fit. At each of L26/L30/L34, we concatenate the commit state, the eighth post-commit state, the mean of the first eight post-commit states, and the running mean before commit (17,920 dimensions per layer). Each layer has a `StandardScaler`→ L_2 -logistic head ($C=10^{-4}$); we normalize the three training logits by their standard deviations and average them. A GPT-5.4-mini judge supplies pass/fail labels against reference answers. We select the recipe on calibration-only LOPO metrics, report five-fold stratified OOF AUC, and score the four external pools without fitting on their labels. Because the variant review repeatedly reported these external metrics, they are development evidence rather than untouched holdouts. The 15/30/50% thresholds are OOF-score quantiles. Because the feature window ends 640 ms after commit, this is an answer-onset read rather than a strictly causal pre-answer read.

Layer location. On the same expanded calibration cohort, the single-feature answer-onset sweep rises from 0.720 OOF AUC at L2 to 0.785 at L30, stays within 0.006 from L18 through L34, and falls to 0.761 at L54. Thus the readable band remains mid-network in this second model family; the reported L26/L30/L34 average spans its plateau.

Scope of the transfer test. The capture follows the model’s native frame protocol, but it is offline: NeMo’s released `offline_voicechat` path re-runs the full prefix every 80 ms and cannot sustain a real-time duplex loop. We therefore use this experiment only to test whether the probe method transfers to another model family and to external speech-QA pools. It does not measure real-time full-duplex serving, relay quality, or end-to-end latency.

Pass-3 results. The three-layer answer-onset probe reaches OOF AUC 0.818, pooled LOPO 0.792, and mean within-held-pool LOPO 0.691. With no external labels used for fitting, its AUC under the GPT-5.4-mini training judge is 0.827/0.862/0.790/0.775 on TriviaQA/WebQ/Llama-Q/SD-QA (mean 0.813; Table 8). Under the pool-aligned judges used for the cross-family comparison, the same scores give 0.845/0.845/0.803/0.770 (mean 0.816). At fixed 15/30/50% per-pool escalation

budgets, four-pool mean accuracy is 0.550/0.650/0.757 versus 0.507/0.577/0.670 for matched-rate random. It beats random on every pool at 30% and 50% ($p \leq 0.0008$); SD-QA at 15% is inconclusive ($p=0.063$). On VoiceBench AlpacaEval, the same pass-3 ranking raises the 1–5 score from 3.55 locally to 3.89/4.17/4.51 at 15/30/50%, versus 3.76/3.96/4.23 for matched-rate random ($n=198$, all $p < .001$).

Read-point boundary. The reported feature window covers the first 640 ms after the model’s commit, so it reads answer onset. A strictly causal probe that uses only the commit state, the preceding eight frames, and the running pre-commit mean loses about 0.03 external AUC (Figure 5). Replacing the full-user-audio mean with the online running mean changes external AUC by only -0.004 $[-0.009, +0.001]$. Fixed calibration thresholds still drift across pools: the latest probe fires on 3/9/20% of Llama Questions and 38/50/73% of SD-QA at nominal 15/30/50% budgets. Per-pool or windowed quantiles are therefore required for budget control.

Table 8: Pass-3 NVDA probe read points under GPT-5.4-mini labels (L_2 -logistic $C=10^{-4}$ on standardized features; 2,481-row analysis fits; external labels excluded from fitting, $n = 249/250/249/200$). $\text{LOPO}_{\text{pool}}$ is mean within-held-pool AUC over 13 eligible calibration pools. Δ is the paired-bootstrap difference in external mean from the single-layer answer-onset reference on the same replay.

read	features	OOF	$\text{LOPO}_{\text{pool}}$	external mean	Δ [95% CI]
strictly causal, L34	commit frame pre-means running mean	0.812	0.694	0.781	-0.029 $[-0.048, -0.009]$
strictly causal, L30	same	0.805	0.678	0.785	-0.025 $[-0.044, -0.006]$
single-layer answer onset, L30	onset-last onset-means user-mean	0.816	0.685	0.809	—
single-layer with running mean, L30	onset-last onset-means running mean	0.814	0.684	0.805	-0.004 $[-0.009, +0.001]$
reported 3-layer answer onset, L26/30/34	+ commit frame, running mean	0.818	0.691	0.813	$+0.004$ $[-0.005, +0.012]$




Figure 5: Pass-3 NVDA method-transfer results under native-frame replay. (a) External AUC under GPT-5.4-mini labels by read point: the reported three-layer answer-onset probe reaches 0.813 mean AUC, while removing post-commit states costs about 0.03. Dots are the four pools. (b) Accuracy at fixed escalation budgets under official judges: the reported probe beats matched-rate random at every budget on average.

767 D What breaks at the output: mechanism audit

768 §3.2 reports that the final-layer last-token read inverts. This section reports what we could and could
 769 not establish about why. Everything below uses the Phase-5d all-layer prefill dumps (every decoder
 770 layer, last-token and mean-pooled, same 600 queries) and the identical probe protocol as the layer
 771 sweep: calibration split, L_2 logistic regression, LOPO trained on the four non-math pools and scored
 772 on held-out math. The reproduction check recovers the published L22 and final-layer numbers to
 773 ≤ 0.012 (scripts/14-17_*.py).

774 **Four mechanisms ruled out.** Table 9 lists them. (i) **Speak mode.** Prefilling the same text queries
775 under the TTS/speak-mode template leaves the cliff exactly where it was (final-layer LOPO math
776 0.362 vs. 0.366 plain), so the readout is not being displaced by a prepare-to-speak state. (ii) **Modality**
777 **axis.** The mean audio-minus-text direction at the final layer is orthogonal to both the inverting final-
778 layer probe and the transferring L22 probe (cosine 0.013 and 0.027; chance is $1/\sqrt{4096} = 0.016$), so
779 the cliff is not the audio/text split that §3.3 measures. (iii) **More shortcut available late.** Source-pool
780 identity is linearly decodable at 0.88–0.96 at every depth, in the duplex model and its raw backbone
781 alike: the query-type shortcut of §3.1 is not something the late layers add; it is what remains once the
782 competence signal stops being readable. (iv) **Wholesale representational drift.** Raw linear CKA
783 between the duplex model and its backbone appears to collapse late (0.80→0.47), but transformer
784 residual streams carry a few massive-activation coordinates that dominate the Gram matrix; with
785 per-feature standardization the same comparison is nearly flat (0.82→0.78) and does not localize
786 the cliff. Projecting out the massive-activation axes, or the pool-mean directions, rescues nothing
787 (0.35–0.40 at every $k \leq 20$, against a random-direction control of 0.37).

Table 9: Mechanism candidates for the final-layer inversion. Final-layer LOPO hard-math AUC on MiniCPM-o 4.5 unless noted; 0.366 is the unmodified baseline and 0.931 the L22 read.

candidate	test	verdict
prepare-to-speak state	speak-mode template prefill	0.362, no change
modality (audio vs. text) axis	cosine with the probe direction	0.013 $\approx$ chance
late layers encode more pool identity	pool decoding by depth	0.88–0.96 at all depths
representation rewritten late	standardized CKA vs. backbone	0.82→0.78, flat
massive-activation coordinates	project out top- k such axes	0.35–0.40, no rescue
query-type directions	project out pool-mean directions	0.34–0.45, no rescue
delivery to the output position	residual component of the read	0.93→0.27 (Fig. 6)

788 **What does separate the models.** Figure 6 measures the probe rather than operating on the model.
789 For the probe trained at each depth on the four non-math pools, we split its held-out score into the
790 part carried by that layer’s five dominant principal directions (estimated from training rows only) and
791 the residual, $s = w^\top Px + w^\top (I - P)x$. The residual component is the distributed read, and it is
792 the one that transfers: it holds 0.85–0.93 through mid-network on MiniCPM-o 4.5 and then decays
793 to **0.27** at the output, with the same decay on MiniCPM-o 2.6 (0.81→0.52) and no decay on either
794 raw backbone (Qwen3-8B 0.86, Qwen2.5-7B 0.71 at their final layers) or on the omni-streaming
795 control (Qwen2.5-Omni holds 0.75 to its output with a flat variance share): the decay tracks duplex
796 fine-tuning specifically, not streaming or audio fine-tuning in general. On MiniCPM-o 4.5 the score
797 also concentrates: the dominant subspace carries 39% of the held-out score variance at the final layer
798 against ≤ 0.18 at every earlier depth and 0.05 for its backbone; on MiniCPM-o 2.6 this concentration
799 is present but weak (0.16 vs. 0.14 for its backbone), so we report it as a property of the stronger
800 duplex model, not a general law.

801 **Direct turn-control tests.** Two experiments target the turn-control story head-on rather than
802 by exclusion (modal_interp.py, scripts/19_*.py). **Control-token logit lens:** applying each
803 model’s own final norm and unembedding to the stored per-layer states, the probability mass on
804 the trained-in control vocabulary (added/special tokens) rises over MiniCPM-o 4.5’s last layers
805 (log-mass −23.8 at L32 to −11.8 at the output) while the raw backbone’s falls by fifteen log-units
806 over the same span: the duplex output layer, uniquely, stays ready to emit control tokens. But
807 the fine-grained prediction fails: the inverting probe direction is no more aligned with the control
808 tokens’ unembedding rows than with random vocabulary rows (max $|\cos|$ 0.069 vs. 0.068), and
809 per-query control mass correlates with the probe score at only $r=0.18$. **Listen/speak intervention:**
810 re-prefilling 60 calibration queries with a still-listening vs. answer-now suffix moves the final-layer
811 state substantially in both models (relative displacement 0.27 duplex, 0.59 raw; cue text is just prompt
812 semantics), and the cue direction is orthogonal to the inverting probe ($|\cos| = 0.019$, chance 0.016),
813 shifting its score by only +0.23 sd (and the deployed L22 read by +0.08 sd). Together: the duplex
814 output layer is measurably control-biased in its output distribution, but the inversion is not carried by

control-token directions or by an explicit listen/speak axis: the simple takeover story fails its own direct tests.

What we do not claim. Destructive tests do not settle the mechanism. Removing the top five principal directions from the final layer does raise LOPO math from 0.378 to 0.758, but the identical operation costs the raw backbone 0.90→0.14 and costs the intact L22 read 0.931→0.760, so it is a damaging intervention that happens to help a broken readout, not evidence of a duplex-added subspace. We therefore state the finding at the level the evidence supports: duplex fine-tuning does not erase the competence signal, it stops delivering it in distributed form to the final last-token position, and the layer where delivery fails is also the one whose output distribution stays control-biased. Attributing the loss to a specific duplex objective needs training-time access we do not have; the practical consequence, read mid-network, does not depend on the attribution.




Figure 6: Non-destructive decomposition of the probe’s held-out score at every depth. **Left:** share of score variance carried by the layer’s five dominant directions. **Right:** transfer AUC of the residual (distributed) component, intact to the output on both raw backbones and on the omni-streaming control, decaying to inversion over the final layers of both duplex models.

E Query-surface router audit

A RouteLLM-style TF-IDF→LR router under identical labels reaches OOF AUC 0.669 (below the pool oracle’s 0.715); a 24-configuration sweep tops out at 0.689, so this is the query surface’s ceiling, not tuning. Under LOPO it collapses to 0.38–0.57 and scores the 100%-fail trap pool at 0.230: it misses the pool a router most needs. The training receipt: at the argmax threshold its out-of-fold accuracy exactly matches the majority class; the internal probes under the same accounting clear majority by 9–29 points (Figure 8). Nor is it data-starved: the identical recipe on RouterBench [Hu et al., 2024] (36,497 prompts) reaches only in-domain AUC 0.710 with leave-one-benchmark-out at chance; RouteLLM’s released preference-trained routers score 0.523/0.533 zero-shot on our pool and rank the trap pool below ordinary hard pools; our router transfers to RouterBench below chance (0.440). On audio labels the router improves (test 0.814) but still trails the mid-layer probe (0.879) and cannot fire mid-stream. In distribution, the gate’s area over random (+0.054) is statistically on par with the pool oracle’s (+0.042; paired bootstrap delta CI [−0.003, +0.027], $p=0.13$): the gate’s case is transfer, not the in-distribution margin. $p(\text{True})$ tradeoff variants: Figure 9.

F Gate timing details

Against the model’s own thinking mode. MiniCPM-o 4.5 ships a built-in reasoning mode, self-escalation in place. Gated-cloud escalation dominates all-THINK on both accuracy and latency at every escalation rate (gated-cloud@33%: 78.7% at mean 6.5 s vs. all-THINK 63.7% at 22.4 s; Table 10). The reason is mechanistic: the gate flags knowledge and trap failures, which extra compute cannot fix, so external escalation is necessary, not merely better.




Figure 7: Accuracy vs. escalation rate, frozen test split: gated escalation vs. random-escalation and pool-oracle (type-only) baselines between the small-only and big-only endpoints.




Figure 8: Identical training receipt for the external router and the two internal probes: the router never clears the majority-class baseline; both probes do, by 9–29 points.

846 **Measurement protocol.** Gate latency is wall-clock from prefill start to the L22 score being available
847 (truncated forward + probe dot product), CUDA-synchronized, one H100, 50 queries per modality,
848 three warmup calls excluded, interleaved per query with the reference arms; P50/P95: gate 20/25 ms
849 text and 45/104 ms audio vs. full-prefill TTFT 36/47 and 68/144 ms. Network and expert-RPC
850 initiation are excluded: the cloud call fires asynchronously at the gate, and its user-facing cost is
851 reported as expert-content-audible latency in Appendix G. Absolute prefill wall-times vary with
852 call overhead across benches; the robust cross-bench statistic is the ratio: the L22 state is ready at
853 $\sim 57\text{--}61\%$ of prefill.

854 Per-layer truncated-forward cost vs. quality (Figure 10): the L22 hidden state is available at $\sim 57\text{--}61\%$
855 of prefill wall-time while L22 is simultaneously the in-mix quality peak; forking much earlier is both
856 weaker and modality-specific, so the fork belongs at $\sim 55\text{--}60\%$ depth. If the cloud call launches at




Figure 9: Offline tradeoff for the $p(\text{True})$ signals (pre-answer and post-draft).

Table 10: Deployment policies, frozen test split ($n=240$); latencies measured.

policy	acc (%)	lat. mean (s)	P50	P95
fast-only	58.8	3.0	3.5	4.2
all-THINK (built-in)	63.7	22.4	17.2	60.1
gated-think @33%	61.3	12.1	3.6	47.6
gated-cloud @15%	68.8	5.3	3.6	10.1
gated-cloud @33%	78.7	6.5	3.6	20.8
gated-cloud @50%	85.8	7.0	3.6	24.4

the gate, the expert result arrives before the local answer finishes for 40–58% of the whole mix but only ~20–31% of escalated queries (they skew to short local answers); the residual wait is P50 2–3 s, P90 ~27 s (Figure 11), the stall-phrase budget the injection protocol must cover.

G Live-loop details

What the signal detects. The pre-answer probe is specifically a retrieval-failure detector, strongest on missing-knowledge and long-tail-trap failures where the empty lookup is observable in the model’s state. Execution failures (stuck mid-derivation) surface in decode-time signals instead, and metacognitive failures (false-premise questions, where no failure event occurs at all) are invisible to every signal we tested, query-surface routers included (escalation rate 18% vs. trap’s 90%; Appendix G), which motivates a decode-time check behind the pre-answer gate.

Full live tables. Table 11 (both scoring views, speakable subset and full pool) and Table 12 (measured response latency). The conservative tier shortens the P95 tail: its few escalations displace exactly the longest local decodes. Figure 12 joins accuracy and latency; Figure 13 replays three real sessions from recorded timers.

Time to first audible audio. The sweep above is text-out; a TTS-on re-run of the balanced arm (same 240 queries, talker head + vocoder on, stall pre-synthesized once in the talker’s own voice)




Figure 10: Fork-depth Pareto: per-layer wall-time vs. gate quality.




Figure 11: Cloud round-trip vs. local answer duration: probability the expert result arrives before the local answer ends.

873 measures speech onset directly. The gate decision is unchanged: the TTS token enters the context
874 only after the end-of-turn read (read p50 21 ms / p99 32 ms), and the same 84/240 queries fire. From
875 gate decision to first audible audio: **local answers p50 0.552 s / p95 0.599 / p99 0.644; escalated**
876 **turns p50 0.022 s / p99 0.029**: the canned stall (3.2–4.1 s of speech) plays the moment the gate fires,
877 so the escalated path reaches first audio 25× faster. Expert content becomes audible at p50 6.4 s /
878 p95 26.0 / p99 34.0 (cache-corrected expert round-trip + first relay PCM, p50 0.66 s after the answer
879 returns); per-turn dead air between stall end and expert content is p50 2.5 s / p95 22.4 / p99 30.7, with
880 25/84 escalated turns fully covered; the stall covers the head of that wait, and Table 12’s completion
881 profile bounds the tail.

882 **Speculative prefetch, measured and dropped.** Framed as speculative escalation (draft–verify at
883 sentence granularity), its acceptance rate (does an expert answer to the partial transcript answer the
884 full question) is 0.07/0.19/0.51 at 25/50/75% of the audio; at the mid-stream gate’s typical fire point
885 6–8 of 10 calls would be wasted. The trap pool’s floor (0.00/0.10/0.27) shares the back-loaded-core
886 property that breaks mid-stream probe reads.

Table 11: Native full-duplex live sweep on the internal test set ($n=240$), bootstrap 95% CIs ($B=5000$); one live arm per tier. Random = matched-rate mix of the never and always arms. The relay-free bound scores the expert’s returned text directly (always arm: 0.765). Top: native sessions with the deployed TTS relay. Bottom, for reference: the retired turn-based harness loop on the same queries (audio input, text output).

tier	esc	delivered	Δ vs. floor	random
never	0%	0.408 [0.35,0.47]	—	0.408
conservative	8%	0.467 [0.40,0.53]	+0.058	0.433
balanced	25%	0.542 [0.48,0.60]	+0.133	0.482
aggressive	52%	0.629 [0.57,0.69]	+0.221	0.561
always (measured)	100%	0.704 [0.65,0.76]	+0.296	0.704
<i>turn-based harness loop (retired), full pool</i>				
never	0%	0.383 [0.32,0.45]	—	0.383
conservative	16%	0.483 [0.42,0.55]	+0.100	0.420
balanced	35%	0.554 [0.49,0.62]	+0.171	0.465
aggressive	61%	0.596 [0.53,0.66]	+0.212	0.526
always (measured)	100%	0.617	+0.234	0.617

Table 12: Time from the end of user speech to the end of the delivered answer (s), native sessions on the internal test set, $n=240$ per arm. Escalated turns include the stall, the wall-clock wait for the expert, and the spoken relay; the P95/P99 tail is dominated by expert calls that hit the 150 s wait cap and by long spoken answers, and is indicative only at this n .

tier	esc	mean	P50	P95	P99	Δ P50
never	0%	5.7	3.6	17.3	21.2	—
conservative	8%	12.0	4.3	22.3	192	+0.7
balanced	25%	15.3	5.6	39.9	290	+2.0
aggressive	52%	27.3	10.9	103	403	+7.3
always	100%	45.1	20.8	210	370	+17.2

887 **Conflicting-injection controls (172 sessions).** Fabricated conflicting expert answers by within-pool
888 derangement, four framings: worst-case comply is 0.00 neutral / 0.53 deployed authority / 0.79
889 strong-override, while seeding words into the talker’s mouth backfires (comply 0.25, lip-service 0.62;
890 forced speech is treated as revisable). Math is silently re-derived rather than relayed or disputed.
891 Model-wrong $\times$ wrong-injection swallow rate is 0.64: expert quality is the accuracy ceiling of the
892 cascade.

893 **Uplink diagnostics.** On the escalated population the expert answering the talker’s own transcript
894 scores 0.585 vs. 0.871 on gold text. Prompt-level rescues on the self-transcript are no-ops (the
895 damaging errors are internally-coherent entity substitutions); replacing the ears helps: open-weights
896 Whisper recovers little (+4 pp) but a stronger cloud ASR on the same audio recovers 38% of the
897 gap (0.585 $\rightarrow$ 0.694, McNemar $p=0.007$), diagnostic only (its critical-path latency and external-pool
898 survival are unmeasured). A same-model gold-text control with an audio-capable expert shows the
899 audio channel itself is lossless for speakable content: the leak is the talker’s transcription, not the
900 synthesis or the relay (relay drops: 5 cases of 108; corrupted transcripts dominate the escalated
901 failures, first-generation sweep).

902 **Boundary query classes. False premises:** the model fails substantially (0.63 text / 0.47 audio)
903 and the expert would help (0.80), but the gate is blind (fire@balanced 0.18 vs. trap’s 0.90; no score
904 separation); answering along a false premise does not feel like inability. Real-time queries (“what
905 is NVDA trading at”): the gate is not blind, the class was simply absent from every static training
906 family; extending calibration with FreshQA [Vu et al., 2023] (fast-changing labeled escalate a priori,
907 never-changing as in-family controls) raises held-out fast-changing escalation rates to 0.80/0.93/1.0
908 across tiers with frozen-test AUC unchanged (0.877 vs. 0.879). One calibration subtlety: tier budgets




Figure 12: Delivered accuracy vs. per-arm P50 time-to-answer, native sessions on the internal test set.

must remain quantiles of the deployment-mix scores, or the new capability is spent on its own threshold shift.

Demonstration system. The pipeline runs as an interactive web demo on one H100, on the model’s native full-duplex head (Appendix O): every second of microphone audio is prefilled into the context the talker is generating from, the head itself decides listen/speak per chunk, and the gate reads L22 at the listen→speak transition against the in-regime thresholds. Escalated turns play a canned stall (teacher-forced once, the talker’s own voice), consult the expert (with a demo-only web-search tool for real-time queries), and the verified answer is spoken verbatim in the talker’s voice, paced on the duplex chunk cadence; whatever was actually played—the stall line, the answer up to any interruption—is teacher-forced into the head’s context as its own completed speech turn, so the context the head takes its next listen/speak decision from matches what the user heard, the same shape a native turn leaves. The delivery is an ordinary duplex turn, hence interruptible like any other. There is no VAD, no ASR barge-in classifier, and no abort machinery: floor control is entirely the duplex head’s own behavior (Appendix O). An earlier demo build approximated barge-in with an energy-VAD + burst-ASR harness over the turn-based loop; it is retired, and survives only as the harness-overlap arm measured above. The demo will be made available upon publication.

H Duplex-validation sweep: the gate under the talker

The talker-disabled live sweep (§5.1) never engages the talker head; this sweep replays its full 240-query pool through the deployed voice loop above to test whether the frozen probe read survives the spoken regime. Two arms, one session per query, deployed v4 gate with scores logged but escalation disabled (the arm isolates gate validity, not routing outcomes). **Clean:** the query streams as mic-shaped PCM (pool TTS audio plus steady noise at RMS 0.008), the energy VAD ends the turn, the talker answers aloud, and the session is cut after the first audible chunk. **Overlap:** an easy warmup question (probe score <0.25, locally answered, >5 s spoken answer) puts the talker mid-answer, and the target query is then spoken over it; sustained overlapping speech commits a barge-in whose audio seeds the next turn. The barge-in path engaged on 237/239 pairs.

Three live sessions, timestamped (balanced arm; every event time is a recorded live-loop timer)




Figure 13: Three real sessions from recorded timers: expert beats the stall phrase (seamless); expert outlives it (dead air, the P99 signature); gate stays closed (local answer fully decoded 0.4 s after end of turn). Speech occupancy is estimated at 150 wpm (hatched); the measured arms are text-out, so no vocoder onset is plotted.

Against the frozen local-failure labels of the same pool, the identical frozen probe reads: talker-disabled AUC 0.866 [0.817, 0.909], clean 0.871 [0.824, 0.914], overlap 0.856 [0.803, 0.901]; barged-only subset ($n=237$) 0.854. Paired bootstrap deltas against the talker-disabled arm are $+0.005$ [$-0.025, +0.035$] (clean) and -0.009 [$-0.036, +0.016$] (overlap), both inside the ± 0.02 – 0.03 live replication floor. Per-query scores correlate across regimes at $r=0.90$ – 0.92 , and the balanced threshold makes the same escalate-or-hold decision as the talker-disabled read on 0.887 (clean) and 0.912 (overlap) of queries. The end-of-turn read costs 24–26 ms at p50 ($p95 \leq 30$ ms) in both arms, with the first audible answer chunk arriving ≥ 0.5 s after the gate decision ($p50 \approx 0.58$ s), so the decision precedes any audible commitment. Scope: this loop keeps per-turn session control: the mic stays open while the talker speaks and interruptions seed the next read, but incoming audio is not prefilled during generation. In hindsight this sweep and the concurrent arm are best read as read-point ablations: the deployed native full-duplex regime (Appendix O) sits strictly between the turn-based ideal and the harness-concurrent worst case on every pool measured. The concurrent regime is measured separately below.

Floor control under escalation: barge-in vs. backchannel. The demonstration loop’s floor policy (speech over the talker ducks playback; a fast ASR pass returns the floor on backchannels or commits the interrupt) raises a distinct worry: did adding the gate change the talker’s overlap behavior? A dedicated sweep injects overlap stimuli into the deployed loop at real-time pacing, four classes: short backchannels (0.40–0.46 s, the English+Chinese lexicon), long lexicon-only continuers (time-stretched to 1.33–2.08 s, probing the sustained-speech commit rule), out-of-lexicon stop commands, and pool queries spoken over the answer, in matched pairs with the gate on and off under an identical floor policy (non-firing queries, $n=40$ per class per arm). The answer is exact orthogonality: all 160 pairs make the identical duck/resume/interrupt decision in both arms (every paired rate delta

0.000 [0.000, 0.000]; paired latency medians within 31 ms), down to the single seeding failure, which occurs on the same query in both arms. By construction the gate reads once at end of turn and never enters the floor state machine; this measures it. The three phases that exist only under escalation (the stall sentence, the expert wait, the relay speech) are the actual new risk surface and behave (Table 13): a genuine barge-in aborts the escalated turn and seeds the next turn 40/40, and short backchannels hold the floor at the same rate as during local answers. One asymmetry is disclosed: a false interrupt during the expert wait cancels the pending expert call, so the policy’s fail-toward-interrupt bias is recoverable but not free there. That residual backchannel false-interrupt rate (0.40 short, 0.50 long; identical in both arms) is a property of the harness policy and is per-stimulus deterministic: non-lexical hums (“mm-hm” and its Chinese counterpart) die by the ASR-empty fail-closed rule (16/16 each; worded tokens such as “okay” 0/16), and continuers sustaining loud speech past the ≥ 1.2 s commit rule die without any ASR check (20/20; pause-bearing variants 0/20). We keep the policy as deployed for measurement integrity and note both rules as the obvious fixes. For scope: the checkpoint’s native duplex profile, measured separately on Full-Duplex-Bench v1.0 [Lin et al., 2025] (727 samples, official scripts), shows best-in-class pause handling (turn-over rate 0.125/0.117 Candor/synthetic), competitive user-interruption response (0.915 at 0.90 s), and zero spontaneous backchannels: the lexicon floor is a harness-level approximation added precisely because the talker head has no native backchannel behavior to preserve.

overlap injected during	backchannel holds floor	barge-in aborts	seeds next turn
local answer (gate off)	0.60 [0.45, 0.75]	40/40	39/40
local answer (gate on)	0.60 [0.45, 0.75]	40/40	39/40
stall sentence	0.63 [0.38, 0.88]	16/16	16/16
expert wait	0.60 [0.33, 0.87]	14/14	14/14
relay speech	0.55 [0.27, 0.82]	10/10	10/10

Table 13: Floor-control sweep over the deployed voice loop. Short backchannels (lexicon bursts) vs. genuine barge-ins (pool queries spoken over the talker), by the phase the overlap lands in; the last three phases exist only under escalation. Interrupt latency p50 1.39–1.40 s from overlap onset in every phase; escalated-turn aborts discard the pending expert answer and the barge-in speech becomes the next turn.

Concurrent prefill: interleaved listen/speak in one KV stream. A third arm drives the streaming API in time division: while the talker is mid-answer, each 1 s chunk of the target query is prefilled into the same session between generation chunks, so the end-of-turn read fires while generation is still active (238/240 queries, full interleave coverage). The carrier answer is teacher-forced: under sampling the duplex model reliably yields the turn: it emits EOS within ~ 2 s of user audio appearing in its context, the trained turn-taking behavior surfacing under interleave. This is the one shift on the transfer ladder that breaks the frozen read: scores drift up (mean 0.48 $\rightarrow$ 0.71; the frozen balanced threshold’s escalation rate jumps 37 $\rightarrow$ 75%), AUC drops 0.869 $\rightarrow$ 0.758 [0.693, 0.816] (paired Δ AUC -0.109 [$-0.158, -0.064$]), and label-free requantiling recovers decision agreement only to 75%. The information, however, survives in the state: refitting the same 12288-d linear read on 360 calibration queries collected in-regime reaches **0.818** [0.761, 0.871] on the held-out 240 (calibration OOF 0.775), more than half the gap recovered from those 360 in-regime rows alone. The mid-layer competence direction persists under concurrent listen/speak; unlike every earlier shift (modality, external pools, talker enabled), calibration must follow the regime. Run-to-run replication of the concurrent read correlates at 0.895. Caveat: the interleave goes through the public streaming API: turn-role headers enter mid-generation and the carrier text is forced, so this measures the concurrent state, not the official duplex serving format.

The full loop in the concurrent regime. With the in-regime gate (thresholds = label-free quantiles of each pool’s own concurrent always-local arm scores), the complete escalation loop (carrier speaking, target audio interleaved, EOT read mid-generation, stall $\rightarrow$ expert $\rightarrow$ relay) was re-run end to end on all seven pools (Table 14). We report this sweep as exploratory: one session per query per arm, an in-regime probe calibrated on 360 internal rows only, teacher-forced carrier (the caveat above).

998 The loop’s mechanics reproduce: realized escalation rates track nominal on the external pools (10.8–
 999 51.2%), and latency holds (EOT read p50 23 ms, canned-stall onset 27–34 ms, local first audio p50
 1000 .655 s vs. .552 turn-based, the price of the longer concurrent context), and on the internal pool, where
 1001 calibration matches the regime, so does selection: accuracy climbs 40.4→52.1→66.2% (speaking
 1002 **44.5→56.0→71.1%**) over the balanced and aggressive tiers, clearing the matched-rate random
 1003 remix of the same never/always outcomes ($p=.017$ and 5×10^{-5} ; speaking .009 and 5×10^{-5}), and
 1004 selective escalation beats always-escalate (66.2 vs. 61.7% at two thirds of the calls; speaking 71.1 vs.
 1005 66.5): the turn-based signature result carries over. The matched-random rows of Table 14 bound any
 1006 stronger reading, in three ways. (i) Gains are not monotone everywhere: the internal conservative
 1007 tier under-fires (2.5% realized vs. 15.8% nominal; the top calibration quantile did not transfer to
 1008 the regime; balanced and aggressive did) and lands flat (40.4→40.0, 44.5→43.6, inside the ± 2 –3
 1009 floor). (ii) Externally the 360-row in-regime probe separates from matched-rate random only in spots
 1010 (Llama @50%: 90.8 vs. 86.5, $p=3\times 10^{-4}$; Reasoning zh @30/50%: $p\leq .0024$; AlpacaEval @15%:
 1011 $p=.031$); at the other external operating points the budget gains are what random escalation would
 1012 buy, consistent with this probe’s external AUC of .57–.67 against the turn-based probe’s .68–.81
 1013 (internally, where calibration matches the distribution, it reads .857 vs. the turn-based .879). We then
 1014 tested the calibration-scale attribution directly: the full 2,310-row calibration mix was re-collected in
 1015 the concurrent state and the same read refit at full scale. Scale is real: every English pool lifts (paired
 1016 against the 360-row probe on the same features: TriviaQA .645→.699, WebQ .641→.713, Llama Q.
 1017 .674→.755, SD-QA .639→.701; mean Δ AUC +.068 [.036, .099]) along a monotone data-scaling
 1018 curve (external mean .625→.689 over 360→2,310 rows; internal .818 unchanged), but it buys only
 1019 about a third of the distance to the turn-based .785–.806, and Reasoning zh, absent from the English
 1020 calibration mix, moves within noise (.615→.577 [−.117, .037]). The residue is the regime itself, not
 1021 calibration scale. Table 14’s loop ran with the 360-row probe and its numbers are unchanged. (iii)
 1022 On Llama Questions selective-vs-always is judge- and noise-sensitive (official judge 90.8 vs. 91.2;
 1023 ours 86.4 vs. 85.6; both deltas inside the floor), so we claim that comparison only on the internal
 1024 pool. Floors move by pool, the regime cost story: Llama $\approx$ unchanged, TriviaQA −16, Reasoning
 1025 zh −18 (the English carrier hurts the zh pool most), WebQ +4.8 under the official judge (within the
 1026 ± 2 –3 replication floor; the larger our-judge shift is alias-matching style sensitivity). AlpacaEval rises
 1027 4.36→4.60 (1–5 scale, always 4.76), but only its conservative tier clears matched random ($p=.031$;
 1028 4.40–4.55 across tiers), so the turn-based honest-negative stands.

1029 **A serendipitous recovery behavior.** An accidental run surfaced one bonus behavior worth report-
 1030 ing: when a long question contains a >1.25 s internal pause, the endpointer fires early and the talker
 1031 starts answering the fragment, but the question’s own continuation then triggers barge-in, and the
 1032 remainder becomes the next turn, where the gate simply re-scores it. Early end-pointing thus degrades
 1033 into an extra turn rather than a lost query, the honest duplex-ish recovery available to a turn-based
 1034 loop.

1035 I Fixed-threshold transfer and the expert-benefit oracle

1036 Both analyses below run in the gold-inject view (escalated rows take the expert’s answer to the gold
 1037 text). The deployed-channel view cannot evaluate a freely moving threshold: the most permissive
 1038 arm escalates 50%, so the other half of each pool has no expert outcome on record, whereas the
 1039 ceiling runs cover every id. Labels follow each pool’s own judge (Table 1); the two judges disagree
 1040 by up to 13 points on Web Questions, so they are never mixed within a row. AUC columns are
 1041 recomputed from the deployed sweep’s stored end-of-turn scores; they differ from Table 1’s AUC
 1042 column (captured in a separate scoring pass) by up to 0.03 externally and 0.04 on the internal pool
 1043 (0.840 here vs. 0.879 there): the passes render and prefill the audio independently, and the probe
 1044 score carries render-to-render jitter. AlpacaEval is excluded: it has no binary correctness label on
 1045 either side, only a 1–5 score whose expert ceiling is degenerate (195/199 at 5).

1046 **One absolute threshold, six pools.** The deployed gate adapts per pool, but only through label-free
 1047 quantiles of its own score distribution. To separate what the probe’s ranking carries from what its

Table 14: The full escalation loop in the concurrent regime (exploratory: one session per arm, 360-row in-regime calibration, teacher-forced carrier), all pools, judge-matched to Table 1 (OAB official for TriviaQA/WebQ/Llama; ours for SD-QA/Reasoning/internal; VoiceBench 1–5 for AlpacaEval). Gate = in-regime probe, per-pool label-free quantile thresholds; tier columns are headed by each pool’s realized rate, and the *rand* row under each pool is the matched-rate random remix of the same never/always outcomes (20k draws; $p = P(\text{random} \geq \text{gate})$, marked * $p < .05$, ** $p < .01$). Accuracies in %; AUC = in-regime probe vs. always-local arm failure, decimals. Script: `scripts/20_conclive_random_check.py`.

pool	never	cons.	bal.	aggr.	always	AUC
TriviaQA	50.4	54.0	64.0	71.6	91.6	0.636
<i>rand</i> (10.8/30.8/46.8%)		54.9	63.1	69.6		
WebQ	61.6	65.2	68.8	71.2	81.6	0.636
<i>rand</i> (15.6/30.4/48.8%)		64.7	67.7	71.4		
Llama Q.	81.6	83.2	84.8	90.8**	91.2	0.672
<i>rand</i> (14.8/29.2/51.2%)		83.0	84.4	86.5		
SD-QA	49.5	55.5	60.5	70.5	89.5	0.636
<i>rand</i> (15.0/29.5/50.5%)		55.5	61.3	69.7		
Reason. zh	40.6	48.5	58.4**	64.4**	80.7	0.570
<i>rand</i> (15.3/30.2/45.0%)		46.8	52.7	58.7		
internal (240)	40.4	40.0	52.1*	66.2**	61.7	—
<i>rand</i> (2.5/36.2/66.2%)		41.0	48.1	54.5		
internal speakable (218)	44.5	43.6	56.0**	71.1**	66.5	0.857
<i>rand</i> (2.3/31.2/62.8%)		45.0	51.4	58.3		
AlpacaEval (1–5)	4.36	4.45*	4.53	4.60	4.76	—
<i>rand</i> (10.1/30.7/47.7%)		4.40	4.48	4.55		

scale carries, we freeze the single absolute threshold the internal system actually deployed ($\tau=0.628$, the balanced tier) and apply that number verbatim everywhere. **Scale does not transfer:** the same τ realizes escalation rates from 2.0% (Reasoning QA) to 48.0% (SD-QA), a 24-fold spread, so a fixed cut cannot hold an escalation budget across pools. **Ranking does:** at whatever rate it fires, fixed- τ escalation beats a random reference at the same realized rate on four of five external pools, paired-bootstrap interval excluding zero (Table 15). The exception is Reasoning QA, where τ fires on 2% of queries and there is almost nothing to measure. This is the concrete case for the label-free quantile step: it buys budget control, not discrimination.

Does the gate find failures the expert can fix? The probe is fit and read against local failure, but what a routing system needs is expert benefit: $y=1$ when the local model is wrong and the expert is right. Re-scoring the same frozen scores against this harder label costs little on the retrieval pools (Speech TriviaQA 0.796→0.782, Llama Questions 0.830→0.813) and more where the expert is itself weak (Web Questions 0.773→0.689 against a 0.864 expert ceiling; internal 0.840→0.732). The gate thus remains a usable benefit predictor out of distribution (0.63–0.81) despite never having seen an expert outcome, and the drop is largest exactly where the paper already reports a flat live curve. On Web Questions the mechanism is direct: 75.0% of local failures are expert-fixable pool-wide, but only 65.5% of the failures inside the gate’s top-scoring 15% are, and that share climbs monotonically as the cut widens (0.655/0.685/0.711 at 15/30/50%): the gate’s most confident failure calls are the ones the expert also misses. A benefit-trained refit closes none of this: with expert outcomes collected for all 2,310 calibration queries (gold-text expert, expert-right rate .853), retraining the same 12,288-d read directly on the benefit label moves internal benefit-AUC .732→.759 (paired +.027 [.004, .052]) and leaves every external pool inside noise (Δ −.029 to +.037, all CIs spanning zero), at no internal cost on the fail label (.840→.839). The benefit gap reflects the label’s dependence on the expert’s own failure surface, not a trainable objective mismatch: the expert-agnostic label leaves nothing on the table. Table 16 places the deployed gate against the matched-budget oracle that escalates true benefit cases first: the gate clears random at every pool and budget, capturing a fifth to three fifths of the random-to-oracle span, and the residual gap is the headroom a decode-time or expert-aware second stage would have to claim.

Table 15: One absolute threshold ($\tau=0.628$, frozen on the internal pool) applied verbatim to every pool; gold-inject view. Δ is fixed- τ accuracy minus a random reference at the same realized rate (paired bootstrap over ids, $B=10^4$, seed 42). AUC columns re-score the same frozen scores against local failure and against expert benefit (y = local wrong $\wedge$ expert right, base rate in the last column). Rates and accuracies in %; AUC in decimals.

pool	n	rate	local	acc	rand	Δ 95% CI	AUC _{fail}	AUC _{benefit}	base
Speech TriviaQA	250	28.4	66.4	83.2	75.0	[+4.4, +12.1]	0.796	0.782	30.8
Speech Web Questions	250	36.0	56.8	72.0	67.5	[+0.8, +8.4]	0.773	0.689	32.4
Llama Questions	250	8.0	83.6	86.8	84.3	[+0.6, +4.6]	0.830	0.813	11.2
SD-QA	200	48.0	51.5	80.5	71.4	[+4.5, +13.6]	0.800	0.765	43.5
Reasoning QA (zh)	202	2.0	58.9	60.9	59.5	[-0.1, +3.5]	0.679	0.634	30.2
internal frozen test	240	35.0	38.3	65.4	57.0	[+3.8, +13.1]	0.840	0.732	53.3

Table 16: Oracle frontier at matched escalation budgets (gold-inject view): random reference / deployed quantile gate / benefit oracle, which escalates the true $y=1$ cases first. The gate clears random everywhere; the oracle column is the remaining headroom. Accuracies in %.

pool	15% budget			30% budget			50% budget		
	rand	gate	oracle	rand	gate	oracle	rand	gate	oracle
Speech TriviaQA	71.0	76.0	81.6	75.5	83.6	96.4	81.6	91.2	97.2
Speech Web Questions	61.2	64.0	72.0	65.7	70.0	86.8	71.6	76.4	89.2
Llama Questions	85.0	88.8	94.8	86.4	91.2	94.8	88.2	93.2	94.8
SD-QA	57.7	63.5	66.5	63.9	74.5	81.5	72.2	82.0	95.0
Reasoning QA (zh)	63.1	67.8	73.8	67.4	72.8	89.1	73.0	76.7	89.1
internal frozen test	46.3	49.2	53.3	54.3	60.8	68.3	65.0	74.2	88.3

J Transfer latency shapes

Matched-rate precision. Precision at a fixed escalation rate is capped at base-fail/rate: on Llama Questions (base fail 0.16) the 50% cap is 0.32 and the deployed probe reaches 0.304, 95% of the cap (recall 0.93). Quoting raw precision would misread the strongest pool as the weakest.

The median left-fold. On easy pools per-arm P50 is non-monotonic (Llama Questions 1.52→1.17→1.60/1.42 s): a verified median-of-a-mixture mechanism, not noise. The probe preferentially escalates the queries whose local decode is longest: on retrieval pools wrong answers hedge at length (decode time and answer length $r=0.97$), so a wrongness-selecting probe strips the slow tail as a side effect (escalated-at-15% ids: always-local arm local P50 1.711 s vs. 1.198 s kept; local accuracy 0.32 vs. 0.73). Controls: the probe does not key on length ($\text{corr}(\text{score}, \text{length}) = +0.05$; ≈ 0 conditioned on wrongness), and on SD-QA, where wrong answers are short, the same probe selects wrongness identically and produces no fold. The mean is monotonic (1.65→2.06 s); kept-local P50 falls 1.52→0.43 s across tiers (Figure 14). On Reasoning QA the shape inverts usefully: the local chain-of-thought is spoken, so every reasoning token enters the response clock, while the expert’s chain is hidden behind a flat ~ 3.7 s round-trip: escalation truncates the latency tail (P90 13.25→10.94 s) rather than fattening it.

AlpacaEval scale caveats. The official 4.8 is an offline chat-mode number (our chat-mode control reproduces it at 4.86) and the chat-vs-live gap (4.86 vs. 3.99) is the duplex loop’s spoken-reply style plus a token cap, so only live arms are comparable to each other. Selected queries score 3.90 on the never arm vs. 3.98 for skipped: no discrimination; escalated rows gain (3.90→4.71) because the expert writes better long-form answers; random picks would harvest the same gain.




Figure 14: Mixture decomposition of Llama Questions latency per arm: kept-local P50 falls 1.52→0.43 s as the probe strips the slowest local decodes while escalated rows pay a flat ~3 s expert round-trip; the arm median (dashed) is their mixture.

1097 K Full-pool and per-family figures

1098 Figures 15 (dual-view) on the full frozen pool, and the complete external-family figure set (Ap-
 1099 pendix L); the noise audit behind the ± 0.02 –0.03 floor is Figure 16.




Figure 15: Local (never-arm) accuracy per pool: native full-duplex sessions vs. the retired turn-based harness loop, same judges. The native floor is 1–8 points lower on five pools and 2.5 points higher on the internal set; bars are bootstrap 95% CIs.

1100 L External-benchmark figure set




Figure 16: Replication-noise audit: same query, same audio, re-run across arms. Kept-local verdicts flip 2.3–18.8% by pool; repeated escalation 0.7–10.6%; implied paired SE 0.009–0.028.




Figure 17: Speech TriviaQA, Speech Web Questions, SD-QA: native full-duplex live arms, delivered accuracy vs. realized escalation rate (left) and vs. P50 time-to-answer (right). Orange dashed on TriviaQA and WebQ: the same gate with the retired talker-steering relay.




Figure 18: Llama Questions, Reasoning QA (zh), VoiceBench AlpacaEval: native full-duplex live arms, delivered accuracy (AlpacaEval: judge score, 1–5) vs. realized escalation rate (left) and vs. P50 time-to-answer (right).

1101 M Judge and self-evaluation prompts

1102 All adequacy labels use the following fixed system prompt (structured outputs with a strict {reason,
1103 adequate} schema; the judge never sees which system produced the answer):

1104 You are a strict, fair grader of a small language model’s answers.
1105 You are given a QUESTION, an optional REFERENCE answer, and a
1106 CANDIDATE answer produced by a small model. Decide whether the
1107 CANDIDATE is ADEQUATE, meaning ALL of: (1) factually correct;
1108 (2) any reasoning shown is valid; (3) it actually and completely
1109 answers the question that was asked. For multiple-choice questions
1110 the candidate must land on the correct option (by letter or by
1111 unambiguously stating the option’s content). A partially correct,
1112 hedged, refused, off-topic, or truncated answer is NOT adequate. If
1113 a REFERENCE is provided, grade against it. Be strict but do not
1114 penalize harmless differences in format or extra correct detail.

1115 The verbalized self-evaluation baselines use one fixed phrasing each; $P(\text{Yes})$ is read from the first
1116 generated token’s distribution:

1117 [pre] I am going to show you a question. Do NOT answer it. Judge
1118 honestly whether you yourself would answer it correctly. Question:
1119 {q} Would you answer this question correctly? Reply with exactly one
1120 word: Yes or No.
1121 [post] Here is a question and a proposed answer. Question: {q}
1122 Proposed answer: {a} Is the proposed answer correct? Reply with
1123 exactly one word: Yes or No.

1124 N Reproducibility

1125 All sampling and splits use a fixed seed (42); the calibration/test split is frozen once and the test split
1126 is touched only by the evaluation scripts. Environment pins: torch 2.8.0; transformers 4.51.0 for
1127 MiniCPM-o 4.5, with per-model pins for comparison backbones; all experiments run on single-H100
1128 instances. Probe artifacts (weights, feature configuration, per-pool label-free thresholds), the judge
1129 prompt above, all trace files, and the per-experiment execution log (including per-run GPU and API
1130 spend, on the order of \$800 total) ship in the project repository, released upon publication.

1131 O The native full-duplex regime

1132 Every result so far reads the probe under a turn-based serving loop (harness-controlled end-of-turn)
1133 or its harness-interleaved concurrent variant (Appendix H). MiniCPM-o 4.5 also ships a native
1134 full-duplex mode: audio prefills in 1 s units into the same context the model generates from, and the
1135 head itself emits ⟨listen⟩ or ⟨speaks⟩, yielding the floor via its turn-end token. This is the deployed
1136 regime of the demonstration system, and this section recalibrates and re-validates the gate inside
1137 it. The read point needs no external end-of-turn detector: the chunk where the head first switches
1138 listen→speak is the commit-to-answer moment, and the probe reads there (the tail- k features then
1139 cover the model’s first ~ 8 answer tokens, a generation-side read the head hands us for free). For 87%
1140 of test queries this commit lands within ± 1 chunk (± 1 s) of the true question end.

1141 **Recalibration.** The full 2,310-row calibration mix was re-collected in this regime (fresh duplex
1142 session per query, features at the deployed read point, labels unchanged) and the same 12,288-d linear
1143 read refit. (The final deployed gate goes one step further and also re-labels: each training row’s label
1144 is the judged outcome of the native answer produced in the same duplex session the features were
1145 read from, under the deployed sampling, rather than a separate deterministic turn-based answer. The
1146 two label sets agree on 80.8% of the 5,228 rows; the native failure rate is higher, .541 vs. .463, mostly
1147 turn-based-correct questions the sampled native answer gets wrong. Ranking is unchanged within
1148 noise, Table 17.) On native test-240 features the in-regime probe reaches AUC **0.830** [0.777, 0.878]

Table 17: Native-regime probe AUC at the deployed listen→speak read point, under the OFFICIAL serving configuration (final deployed system; the chronological narrative above measured under inherited wrapper defaults). The merged refit adds the 2,700-row public-benchmark expansion (incl. 400 Mandarin rows) to the official core. En-4 mean .770 vs. turn-based .771: the English regime cost is eliminated; the Mandarin slice restores Reasoning-zh from the all-English chance collapse. The last column is the deployed gate: same 5,228 features, but every training label is the judged outcome of the same native-duplex answer the features were read from (instead of a separate deterministic turn-based answer). Label source is a wash for ranking (En-4 mean .771 either way; the paired bootstrap crosses zero on six of seven pools, TriviaQA +.035 [+ .006, +.065] being the exception), so the deployed gate uses the in-regime definition. Internal test is scored against both the earlier harness labels and the native-judged labels of the same 240 queries.

pool	<i>n</i>	old cfg	official (2,542)	+exp3 (5,228)	+native lbl. (deployed)
internal (harness lbl.)	239	.830	.846	.844	.843
internal (native lbl.)	239	—	—	.878	.876
TriviaQA	250	.711	.759	.767	.802
WebQ	241	.736	.754	.772	.777
Llama-QA	245	.757	.789	.815	.807
SD-QA	200	.736	.737	.728	.700
Reasoning-zh	202	.606	.495	.605	.621

(transferred concurrent probes: 0.807–0.818; turn-based reference 0.877). Externally the in-regime refit beats both transferred probes on all five pools: TriviaQA .711 ($\Delta +.080$ [+ .030, +.129]), WebQ .736, Llama-QA .757 ($\Delta +.118$), SD-QA .736, Reasoning-zh .606 ($\Delta +.115$); external mean .709. The regime ordering is monotone everywhere: turn-based (.877/.771 internal/external mean) > native (.830/.709) > concurrent (.689 external), i.e. a substantial share of the concurrent regime’s cost was the harness read point, not liveness itself. Thresholds again do not transfer (balanced quantile 0.386→0.645); scores are decoding-temperature invariant (0.7-trained probe on temp-0.1 features: AUC 0.843).

Benchmark-overlap audit. The calibration expansions and the external speech benchmarks both draw on public QA sources (TriviaQA appears on both sides of the ledger), so we audited every calibration query against the five external pools: normalized exact match plus a token-Jaccard ≥ 0.8 near-duplicate screen. Eight collisions exist — 4/250 Speech TriviaQA and 4/250 Llama Questions queries have a (near-)duplicate among the deployed 5,228 calibration rows; the other three pools have none. Refitting the deployed gate with the seven colliding calibration rows removed moves external AUC by at most .002 (TriviaQA .811 → .811, Llama Questions .743 → .741 on the official-configuration dumps) and re-mixed delivered accuracy at the deployed tiers by less than half a point, so no transfer result rests on memorized test questions (`scripts/41_leak_audit.py`).

Validity. Re-mixing native local outcomes (gpt-judged, deployed decoding) with cached always-arm expert outcomes under the native scores (Table 18): under the deployed gate, gated accuracy clears matched-rate random on all five English pools at both deployed tiers: frozen .431→.669 at 49% escalation ($p < .0001$; ceiling .669), TriviaQA .568→.852 at 54%, WebQ .411→.680, Llama-QA .767→.841, SD-QA .455→.775 at 61%. In-regime training renormalizes the score distribution, so realized escalation rates sit near their nominal quantiles on the English pools (the earlier 64–73% drift was a transferred-threshold artifact). Reasoning-zh is an operating-point problem before it is a ranking problem: a single global threshold, quantiled on a 92%-English training mix, places the zh score distribution below every tier (1–2% fire at “balanced”), the language axis of Appendix H expressed as a fire-rate failure. The deployed fix needs no evaluation data: tier thresholds are quantiled on the training-set out-of-fold scores per language (en/zh), and the session selects the language’s thresholds. The zh thresholds derived from the 400 Mandarin training rows alone land Reasoning-zh at 17/34/55% fire (nominal 15/30/50) and lift it from .525 to .609/.658/.713, beating matched-rate random at the two lower tiers ($p = .0065$, $p = .023$). This static per-language point is only as good as the match between the language slice’s family mix and the deployment stream: when the Mandarin training slice is widened from 400 to 1,900 rows with higher-failure families (below), its balanced

Table 18: Native-regime gated accuracy under the official serving config and the deployed merged (official+exp3) gate (8ad remix: native scores select, native local outcomes + cached always-arm expert outcomes mix). p = permutation vs. matched-rate random; escalation rate in parentheses. Thresholds are the deployed per-language operating points: tier quantiles of the training-set out-of-fold scores, computed separately for the English and the Mandarin training rows (no evaluation pool is consulted); the Mandarin pool uses the zh thresholds. On frozen test the aggressive tier reaches the always-escalate ceiling at 49% of the expert cost. Reasoning-zh, which the global English thresholds never reached (1–2% escalation), now escalates at its nominal rates and beats matched-rate random at the two lower tiers. AlpacaEval (measured under the 2,542-row official gate) stays a clean honest negative.

pool	local	balanced	aggressive	always
frozen test	.431	.565 (27%, $p=.0005$)	.669 (49%, $p<.0001$)	.669
TriviaQA	.568	.684 (20%, $p=.0005$)	.852 (54%, $p<.0001$)	.916
WebQ	.411	.527 (22%, $p=.045$)	.680 (58%, $p=.020$)	.809
Llama-QA	.767	.792 (4%, $p=.002$)	.841 (16%, $p<.0001$)	.910
SD-QA	.455	.605 (24%, $p=.0035$)	.775 (61%, $p=.003$)	.895
Reasoning-zh (zh thr.)	.525	.658 (34%, $p=.023$)	.713 (55%, $p=.066$)	.807
AlpacaEval (VB 1–5)	3.65	3.70 (6%, n.s.)	4.07 (37%, $p=.24$)	4.76

1182 quantile moves from .39 to .75 and fires 5% on Reasoning-zh; a family-balanced quantile recovers
1183 13/24/38%. The composition-free alternative (each pool’s own score quantile, or its online form, a
1184 100-turn sliding-window tracker with no labels and no pool identity) realizes nominal rates on every
1185 pool and converges within $\pm .06$ of nominal on the English pools.

1186 **Is the Mandarin gap a language gap?** No: it is a task-type gap. A targeted 1,500-row Mandarin
1187 expansion (Chinese SimpleQA long-tail facts, C-Eval and CMMLU knowledge MC, the remaining
1188 XCOPA-zh; public benchmarks, deduplicated against every pool, judged on the native answers; native
1189 failure .578) does not move Reasoning-zh when added to the calibration mix (.621→.627, paired
1190 +.006 [−.026, +.040]; internal test −.010, English external +.006), and Mandarin-only training
1191 is worse everywhere (Reasoning-zh .591, internal .758). Read the other way, the deployed gate
1192 scores these 1,500 unseen Mandarin questions at AUC **.811** [.789, .833] (long-tail .719, C-Eval .730,
1193 CMMLU .642, causal .693), the same band as the English pools. What the probe does not read is
1194 multi-step reasoning, which is what Reasoning-QA is; the “.62 on Mandarin” number is that gap
1195 wearing a language label. The expansion is therefore not deployed and serves as a second, larger
1196 Mandarin held-out pool. (A side observation from the same dump: on the causal-commonsense family
1197 34% of the native answers switch to English while remaining judged adequate, a language-consistency
1198 defect the adequacy label does not see.)

1199 **Reasoning coverage and in-context rows: the read point saturates.** Two further tar-
1200 geted sets test the two remaining hypotheses. A bilingual multi-step reasoning pool (1,319
1201 rows: LogiQA and LogiQA-zh, MATH levels 4–5 without LaTeX, StrategyQA, BBH
1202 date/deduction/causal/temporal/navigation, CMATH grades 4–6, Ape210K; native failure .595) leaves
1203 Reasoning-zh unchanged when added to calibration (.621→.620, paired −.002 [−.039, +.038]) and
1204 moves nothing else beyond noise. The reason is visible inside the pool itself: cross-validated AUC on
1205 the reasoning families is only .61–.76 (overall .736 [.708, .762]), against .77–.82 for fact, knowledge-
1206 MC, and causal families. Whether a multi-step chain will come out right is weakly encoded in
1207 the listen→speak commit state, which precedes the chain; the .62 on Reasoning-zh is that ceiling,
1208 not a coverage gap. Second, the calibration questions re-collected as the second turn of a session
1209 (a spoken carrier question and its answer first) fail more often (.517→.620; 26% of labels flip)
1210 while the deployed gate scores them lower (mean shift −.097): at the balanced tier the in-context
1211 escalation rate halves (.252→.115) although ranking holds (AUC .861 [.815, .896]). Adding these
1212 rows to training does not change AUC either; the defect is the operating point, and the deployed
1213 demo therefore thresholds on a sliding-window quantile of the session’s own recent onset scores
1214 (100-score window, static per-language threshold until 20 scores exist), which follows such shifts

without labels or pool identity. After a label-definition fix and three targeted expansions ($\approx 4,000$ rows), external ranking sits where the 5,228-row gate left it. The head is not the bottleneck either: on the same features, per-dimension standardization costs .07 of English external AUC (the raw activation scale is informative), an elastic-net head is a wash (internal $+.008$, external $-.007$), and non-linear heads are clearly worse (256-unit MLP .702, PCA-256 + gradient-boosted trees .715, vs. .771). Nor are several heads: a three-way task-type router trained on the same features is 97% accurate, yet no family-specialized head beats the single head on its own family (facts .846 vs. .847, reasoning .761 vs. .766, chat .672 vs. .720), and the routed mixture is flat-to-negative externally (SD-QA $-.033$ [$-.060, -.007$]). What the commit-chunk state encodes about the coming answer is extracted by one linear read; the remaining headroom lies in what is read: later into the answer, other layers, or a second signal family such as verbalized self-evaluation (Table 7), not in how. The talker’s own local floor is lower in-regime (frozen .483 $\rightarrow$.431 under the official serving configuration; the inherited-defaults harness measured .371, and the $\approx$ one-third-temperature / ≈ 8 -point decomposition belongs to that older figure), which raises the gate’s marginal value in deployment. The in-regime calibration scaling curve (measured under the default serving configuration) mirrors the concurrent one and is not saturated: internal .717 $\rightarrow$.830, external mean .643 $\rightarrow$.709 from 360 $\rightarrow$ 2,310, and a further public-benchmark expansion to 5,009 rows (2,300 English across the same seven families + 400 Mandarin: MGSM-zh and XCOPA-zh through the same TTS pipeline, source-disjoint from Reasoning-zh) continues the $\approx +.02$ -per-doubling external trend to **.736** while internal stays flat (.834), and the Mandarin slice moves the axis English data never touched, Reasoning-zh .606 $\rightarrow$.659 (Δ vs. conc-2310 $+.169$ [$+.085, +.251$]; the deployed official-configuration merge reads the same pool at .605 with turn-based labels and .621 with the deployed native labels, Table 17). Two caveats on reading this curve: its six points are the project’s chronological accumulation, each step adding new query families, coverage, or a serving-configuration change, not repeated random subsamples of the final set, so it evidences that added coverage transfers, not that another random doubling of the same families would return $+.02$; and the internal/English-external plateau against the turn-based reference (.771) is an empirical baseline at $n \approx 200$ –250 per pool, not a ceiling.

Ceiling, stability, and the pre-answer read. Three diagnostics on the deployed 5,228-row gate. (i) **Ceiling.** A pool-prior scorer (each query scored by its training pool’s out-of-fold failure rate) reaches AUC .786 against the probe’s .848 on the training out-of-fold scores, and .767 against .876 on frozen test; AUC restricted to same-pool pairs is .743 (train) and .820 (test). Per-question discrimination is thus $+.06$ –.11 of AUC on top of the type prior, the same decomposition as the $n=600$ audit of Section 3.1; at a 30% budget the prior alone recovers .475 failure recall against the probe’s .506 (random .299). The external pools, each a single source, measure the within-pool component directly (.70–.81, Table 17). (ii) **Stability.** An independent second native pass over frozen test (deployed sampling, $T=0.7$, its own judged label) changes 98% of answer texts and flips 14% of native labels, yet the read at the commit chunk barely moves: feature cosine .95, score correlation .95, and the tier decisions agree on 91–94% of queries. Each run’s features predict the other run’s outcome about as well as their own (.848/.887 vs. .876/.872), so the probe reads a property of the question in context rather than the fate of one sampled answer; against the 204 queries whose two outcomes agree, AUC is .91–.93, which bounds the single-sample label noise in the reported numbers. (iii) **Pre-answer read.** The same layer-22 read taken before the onset chunk’s generate call (no answer tokens in the tail), scored by the deployed head without refitting, gives AUC .861/.837 against the two label sets, .011–.015 below the deployed read; the first ≈ 8 answer tokens are worth about one AUC point, and the pre-answer variant needs its own operating points (the deployed quantiles fire 0/1/20% on it). A hard-example check closes this route from the training side: at fixed regularization, upweighting the 711 calibration failures the out-of-fold probe misses at the widest tier drops external transfer from mean AUC .743 to .715 (4 $\times$), below every draw in 100-subset random-failure null distributions, count-, weight-, and source-composition-matched (null means .737/.740, empirical $p < .01$) — the harm is specific to those rows, not to extra positive mass. Retuning the regularizer per arm recovers most of the loss (.735 at 4 $\times$, with C' shrinking $3 \times 10^{-4} \rightarrow 3 \times 10^{-5}$ at 8 $\times$): the solver’s best response is to suppress the fit to the upweighted rows, and no weighting improves on the unweighted baseline. The missed failures are therefore not recoverable by emphasis at this read

point — consistent with the execution-slip class carrying no linearly readable signal at this read point (timestamp-limited predictability, not a claim of irreducibility in principle). This is the signature of emphasis on feature-indistinguishable positives rather than under-represented ones [Swayamdipta et al., 2020, Pleiss et al., 2020], and matches the independent finding that internal-state probes read parametric-knowledge recall rather than output truthfulness, leaving association-driven confident errors undetectable [Cheang et al., 2025]; correctness of a reasoning chain becomes linearly readable only as generation proceeds [David, 2025], which is what the read-point sweep above measures in the duplex regime.

Live validation. Full live runs on frozen test-240 (real raw-audio ASR uplink, real expert, wait paced at one chunk per second) validate the remix arithmetic where the channel is thin and expose where it is not: balanced delivers .504 at 23% fire (remix predicted .492 at 24%); aggressive delivers .537 at 45% against a .596 remix prediction, and the gap decomposes entirely onto fired turns: live fired-turn accuracy .495–.589 vs. the .667 cached-expert ceiling. The native relay channel (the answer re-entering as a text unit and being voiced) is the regime’s lossy element (89% of relay units need the nudge retry), playing the role the transcription uplink played in the turn-based loop; the uplink itself is now clean (raw-audio ASR, p50 0.9 s). Latency, expert-cache warm: expert p50 4.4–5.7 s, wait p50 5–6 chunks, fire→relay first audio p50 ~7–8 s. Selection is visible live: unfired subsets score .478/.573 against the .371 unconditional local floor.

Dialogue-act gating. Full duplex creates commits that are not answers: after a barge-in or a backchannel the head speaks to manage the floor (“Okay.”), and the failure probe is out-of-distribution there, on 194 TTS’d floor-management stims (stop commands, backchannels, acknowledgments, fillers; en+zh) the unmodified gate false-fires on 37% at the aggressive tier (stop commands 45%), i.e. it would escalate “stop” to the expert. The fix is a second linear head on the same L22 read, info-seeking vs. floor-management, which separates the two acts perfectly (OOF AUC 1.000); thresholded at the question-side 0.5th percentile it costs 0.5% of true escalations and admits 0% of floor stims. Escalation requires both heads: $\text{act} \geq \tau_a \wedge P(\text{fail}) \geq \tau$. The deployed demo answers floor turns naturally and reports the act score per commit; the failure probe’s calibration is untouched.

Live use then sharpened both heads within minutes: scripted pools miss what ten minutes of real conversation finds. (a) Live mic speech shifts both act distributions into the calibration gap (real request-phrased questions score ~.93, live floor turns ~.37 vs. a standalone-stim maximum of .05): a percentile-anchored τ_a rejected a real question and the talker delivered an empty promise. The fix is a gap-center threshold plus in-context calibration stims, each utterance dumped as the second turn of a carrier conversation, matching how floor acts actually arrive; the act read stays perfectly separable (AUC 1.000) under both conditions, and in-context stop commands false-fire the failure probe on 57% at the aggressive tier, so the act term is not optional. (b) The native refit had silently dropped the FreshQA real-time extension (“the stock price of NVDA today” scored $P(\text{fail})=.29$ and stayed local); replaying that recipe in-regime restores fast-changing heldout fire to .45/.75/1.00 across tiers at zero internal-AUC cost (.830→.833), with budgets still quantiled on the core mix. (c) A relay’s own speak-onset must not re-enter the gate (re-fire with an empty uplink); deliveries are guarded to their end of turn. The pattern across all three: the mid-layer signals separate cleanly every time, and what re-opens at each regime or domain shift is calibration coverage and thresholds: interactive use is itself an evaluation modality.

Multi-turn grounding. A stateless expert uplink breaks on follow-ups: “what is NVDA trading at” then “what about Apple” sends only the second utterance, unresolvable in isolation. The probe itself is unaffected: it reads L22 with the prior turns live in the model’s KV cache, so the fix is in the uplink, not the gate: the deployed demo threads a rolling resolved-dialogue history into the expert input and asks it to resolve references, after which the follow-up correctly returns Apple’s share price. The measured eval pools are single-turn, so this is a deployment capability rather than a table; systematic multi-turn evaluation is future work.

1317 **Floor control, natively.** With the harness gone, 8ba’s question re-poses on the head itself (108
 1318 scripted overlap trials, stims spoken over the talker at fixed offsets). During answer speech, backchan-
 1319 nels almost never stop the floor (false-stop 3/24 within a 6-chunk window) with no ASR or lexicon
 1320 anywhere in the loop; sustained question speech yields the floor (6/12 within 6 chunks, median 8.5
 1321 post-stim); short burst commands (“Stop!”) mostly do not (3/12): the native head reads sustained
 1322 speech, not commands, a naturalness-for-reliability trade the old energy-VAD did not make. The
 1323 escalation machinery is floor-transparent where it should be: relays complete under overlapping
 1324 stims (26/30), and backchannels during the thinker wait leave the pending relay intact (8/10). The
 1325 weak stop-command sensitivity decomposed into two causes, found in sequence. **Mechanism:** in-
 1326 strumenting the serving loop shows the head never samples ⟨listen⟩ mid-turn (zero attempts across
 1327 36 trials), so its only stop channel is the turn-end token. **Configuration:** a user-demanded vanilla
 1328 control (the bare official serving stack, no gate machinery) revealed that our loop had inherited
 1329 the checkpoint wrapper’s sampling defaults ($k=100$) rather than the official demo’s serving config
 1330 ($k=20$, a 3-chunk startup listen guard, an assistant-style system prompt). Under the official config
 1331 the turn-end channel is responsive: a spoken “stop” two seconds into an enumeration answer ends
 1332 the turn in 2.0–2.2 s in 3/3 scripted sessions, versus 1–13 s (once running the answer to completion)
 1333 under the inherited defaults: wide- k sampling dilutes away the rising turn-end probability. The
 1334 deployed demo now serves the official config, and the earlier floor-control stop numbers measured
 1335 under the inherited defaults should be read as a lower bound on native responsiveness. Two durable
 1336 lessons: serving-parameter parity with the reference stack is part of “native”, and a vanilla control
 1337 arm belongs in the harness from day one. The one true interaction: a new question during the wait
 1338 derails the pending relay 7/10 times: the head simply moves on, which is the empirical case for
 1339 aborting the in-flight expert call on user speech, left as deployment future work. Latency: listen
 1340 chunks cost 0.04 s and speak chunks ~ 0.5 s on one H100 (realtime holds); fire→stall-audio 0.2–1 s;
 1341 fire→relay-first-audio is the expert wall clock (13–20 s observed) plus ~ 2.5 s.